

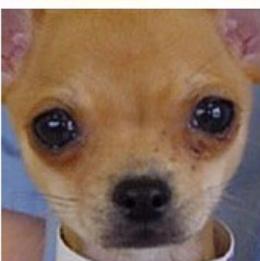
# Convolutional Neural Networks

Slides by Prof. Ben Leong

Let's take a  
step back. . . .

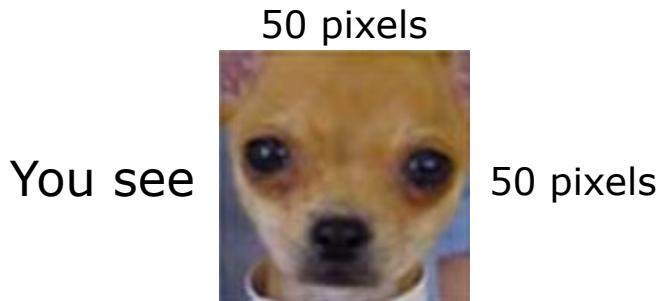
# How do we solve AI problems?

1. Formulate the problem
  - Data/state representation
  - Goal state **Data has structure**  
**Spatial vs Temporal**
2. Apply some known algorithm to well-defined problem





Credit: Internet meme  
original source unknown



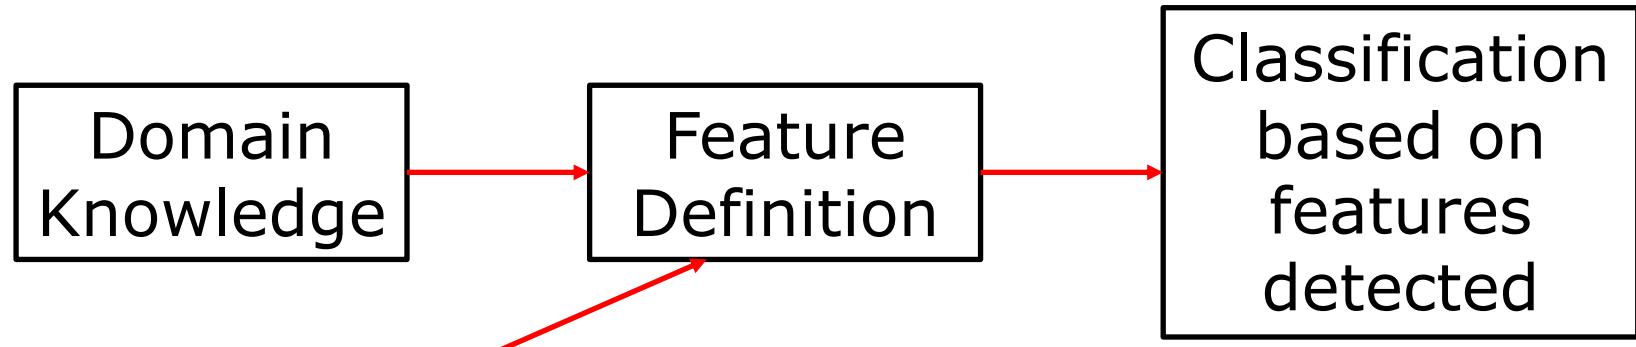
# Chihuahua or muffin?

Computer sees:  $\begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{2,500} \end{bmatrix} = \begin{bmatrix} \text{colour of pixel 0} \\ \text{colour of pixel 1} \\ \vdots \\ \text{colour of pixel 2,500} \end{bmatrix}$

## Question of the Day:

Suppose we restrict terms to max degree 2, how many features do we have?

# “Traditional” approach



**Difficult!**

Scaling &  
Rotations

Deformation  
Occlusion

Illumination  
/shadows

**“Automated” Feature Extraction**

# Naïve approach

Using pixel values directly as an input vector:

- Lose information on spatial structure
- Huge number of weights for the input layer

**Exploit Spatial Structure**

# Images as 2D matrices

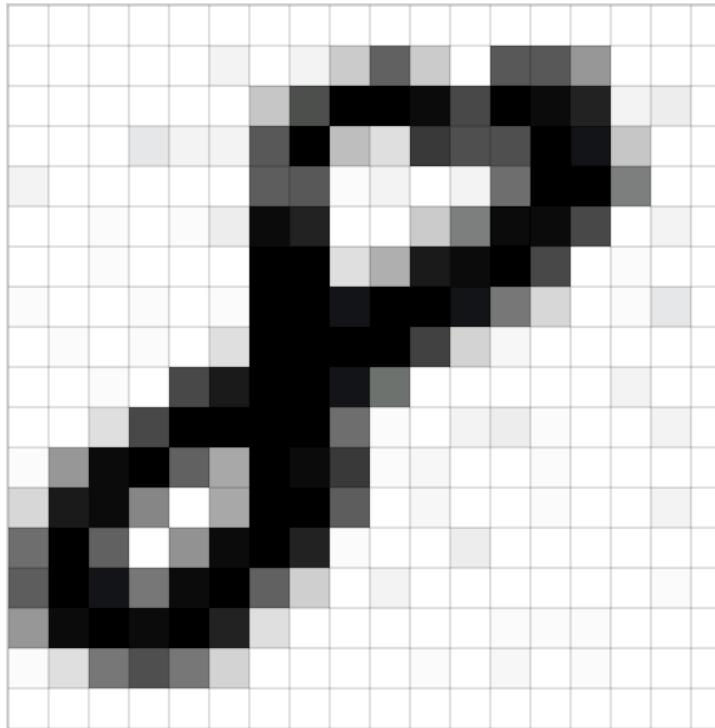
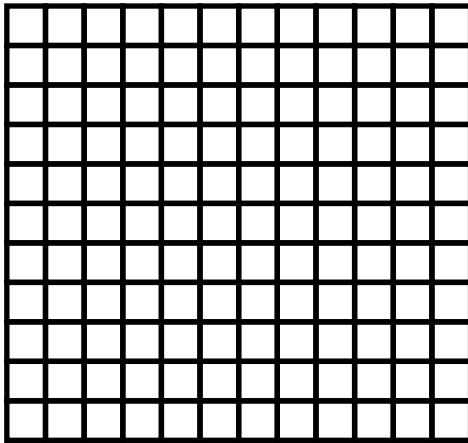


Image credit: <https://medium.com/@ageitgey/machine-learning-is-fun-part-3-deep-learning-and-convolutional-neural-networks-f40359318721>

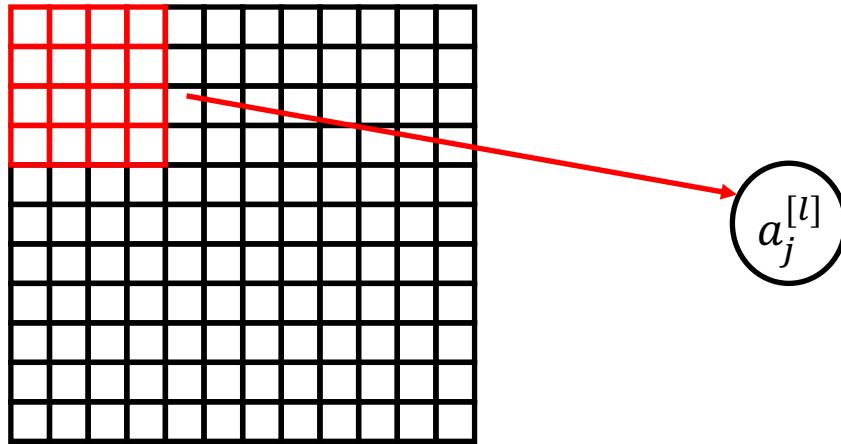
# Convolution

Key Idea: Consider not one pixel at a time, but a group of pixels



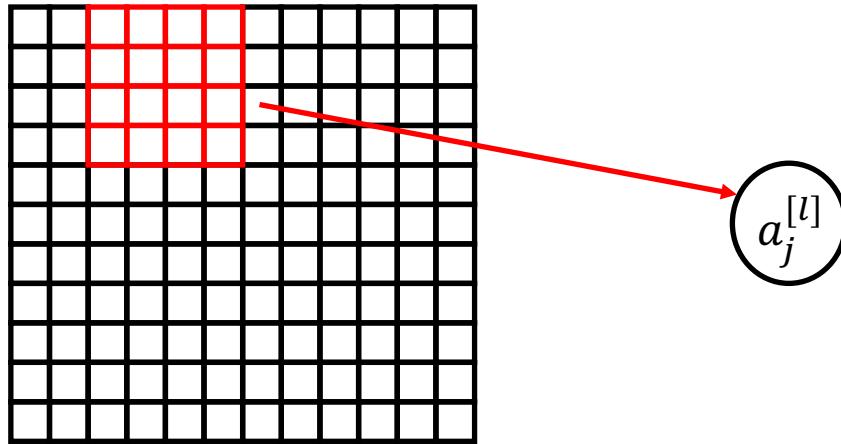
# Convolution

Key Idea: Consider not one pixel at a time, but a group of pixels



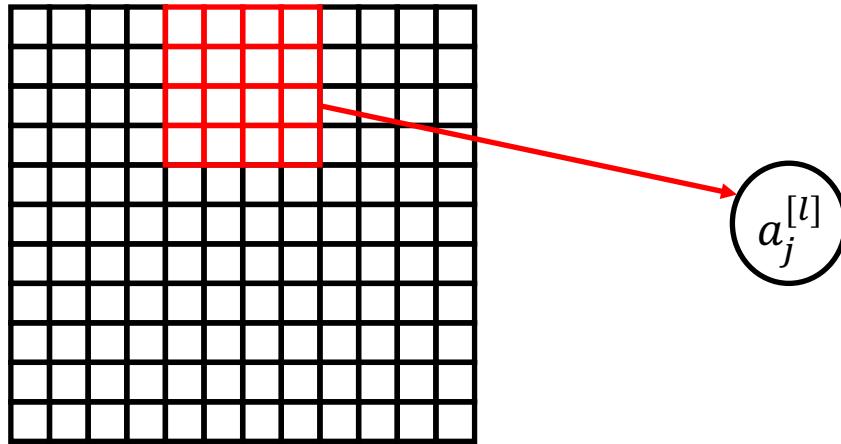
# Convolution

Key Idea: Consider not one pixel at a time, but a group of pixels



# Convolution

Key Idea: Consider not one pixel at a time, but a group of pixels



# How does convolution work?

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image

$$\begin{matrix} & \begin{matrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{matrix} & = & \begin{matrix} 3 \\ \vdots \\ \vdots \end{matrix} \end{matrix}$$

kernel

feature map

Multiply and sum

# How does convolution work?

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image

$$\begin{matrix} & \begin{matrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{matrix} & = & \begin{matrix} 3 & 3 \\ \vdots & \vdots \end{matrix} \end{matrix}$$

kernel

feature map

Multiply and sum

# How does convolution work?

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image

$$\begin{matrix} & \begin{matrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{matrix} & = & \begin{matrix} 3 & 3 & -2 \\ \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots \end{matrix} \end{matrix}$$

kernel

feature map

Multiply and sum

# How does convolution work?

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image

$$\begin{matrix} & \begin{matrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{matrix} & = & \begin{matrix} 3 & 3 & -2 \\ -3 & & \\ & & \end{matrix} \\ \begin{matrix} * & \text{kernel} & \text{feature} \\ & & \text{map} \end{matrix} & & & \end{matrix}$$

Multiply and sum

# How does convolution work?

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image

$$\begin{matrix} & \begin{matrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{matrix} & = & \begin{matrix} 3 & 3 & -2 \\ -3 & -3 & \end{matrix} \\ \begin{matrix} * & \end{matrix} & \text{kernel} & & \text{feature map} \end{matrix}$$

Multiply and sum

# How does convolution work?

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image

$$\begin{matrix} & \begin{matrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{matrix} & = & \begin{matrix} 3 & 3 & -2 \\ -3 & -3 & 4 \\ \end{matrix} \\ \begin{matrix} * & \text{kernel} & \text{feature} \\ & & \text{map} \end{matrix} & & & \end{matrix}$$

Multiply and sum

# How does convolution work?

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image

$$\begin{matrix} & \begin{matrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{matrix} & = & \begin{matrix} 3 & 3 & -2 \\ -3 & -3 & 4 \\ 3 & & \end{matrix} \\ \begin{matrix} * & \text{kernel} & \text{feature} \\ & & \text{map} \end{matrix} & & & \end{matrix}$$

Multiply and sum

# How does convolution work?

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image

$$\begin{matrix} & \begin{matrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{matrix} & = & \begin{matrix} 3 & 3 & -2 \\ -3 & -3 & 4 \\ 3 & 3 & \end{matrix} \\ \begin{matrix} * & \text{kernel} & \text{feature} \\ & & \text{map} \end{matrix} & & & \end{matrix}$$

Multiply and sum

# How does convolution work?

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image

$$\begin{matrix} & \begin{matrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{matrix} & = & \begin{matrix} 3 & 3 & -2 \\ -3 & -3 & 4 \\ 3 & 3 & -4 \end{matrix} \\ \begin{matrix} * & \text{kernel} & \text{feature} \\ & & \text{map} \end{matrix} \end{matrix}$$

Multiply and sum

# How does convolution work?

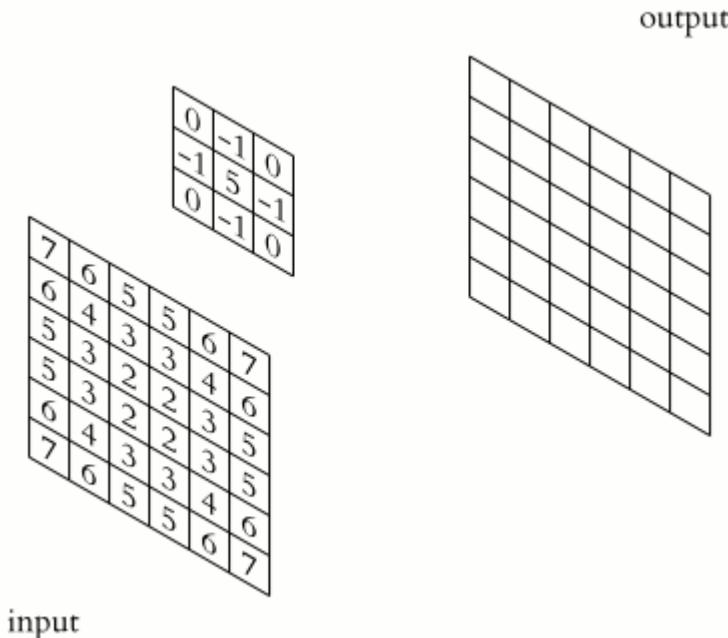


Image credit: <https://towardsdatascience.com/types-of-convolution-kernels-simplified-f040cb307c37>

# Clarification: Convolution

$$(f * g)(t) = \int_{-\infty}^{\infty} f(x)g(t - x)dx$$

Discrete form (2D):

$$x[m, n] * h[m, n] = \sum_{j=-\infty}^{\infty} \sum_{i=-\infty}^{\infty} x[i, j] * h[m - i, n - j]$$

# How does convolution work?

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image

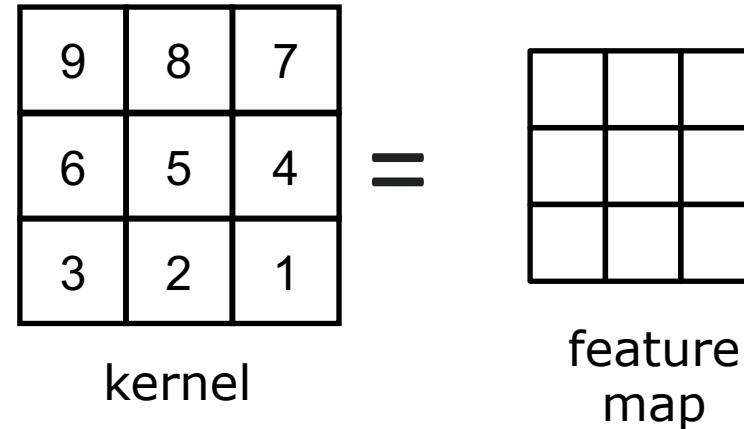
$$\begin{matrix} & \begin{matrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{matrix} & = & \begin{matrix} \square & \square & \square \\ \square & \square & \square \\ \square & \square & \square \end{matrix} \\ \ast & \text{kernel} & & \text{feature map} \end{matrix}$$

Flip first

# How does convolution work?

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image



Flip first  
Multiply and sum

# How does convolution work?

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image

$$x[m, n] * h[m, n] = \sum_{j=-\infty}^{\infty} \sum_{i=-\infty}^{\infty} x[i, j] * h[m + i, n + j]$$

1	2	3
4	5	6
7	8	9

=


feature map

kernel

If we don't flip first  
⇒ 2D Correlation

# Convolution Filters



original



sharpen

0	-1	0
-1	5	-1
0	-1	0



Laplacian

0	1	0
1	-4	1
0	1	0



emboss

-2	-1	0
-1	1	1
0	1	2

# Convolution Filters



original



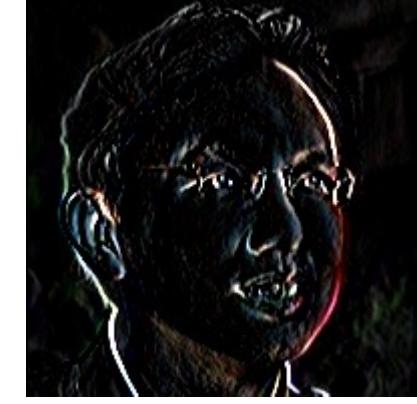
outline

-1	-1	-1
-1	8	-1
-1	-1	-1



top sobel

1	2	1
0	0	0
-1	-2	-1



left sobel

1	0	-1
2	0	-2
1	0	-1

Different filters can extract different features!

# Padding & Stride

- We lose pixels at the edge of our image  $\Rightarrow$ padding
- Computing adjacent filters might be slow  
 $\Rightarrow$ stride

# Padding & Stride

0	1	1	0	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	1	1	1

image

$$\begin{matrix} & \begin{matrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{matrix} & = & \begin{matrix} 3 & 3 & -2 \\ -3 & -3 & 4 \\ 3 & 3 & -4 \end{matrix} \\ \begin{matrix} * \end{matrix} & \text{kernel} & & \text{feature map} \end{matrix}$$

# Padding & Stride

0	0	0	0	0	0	0
0	0	1	1	0	0	0
0	0	1	1	0	0	0
0	1	0	0	1	1	0
0	0	1	1	0	1	0
0	1	1	1	1	1	0
0	0	0	0	0	0	0

image

\*

0	-1	0
-1	5	-1
0	-1	0

kernel

=

-1	3	3	-1	0
-2	3	3	-2	-1
5	-3	-3	4	3
-3	3	3	-4	3
5	2	2	3	3

feature  
map

+ padding  $p_h$  and  $p_w$

# Padding & Stride

0	0	0	0	0	0	0	0
0	0	1	1	0	0	0	0
0	0	1	1	0	0	0	0
0	1	0	0	1	1	1	0
0	0	1	1	0	1	0	0
0	1	1	1	1	1	1	0
0	0	0	0	0	0	0	0

image

\*

0	-1	0
-1	5	-1
0	-1	0

kernel

=

-1	3	3	-1	0
-2	3	3	-2	-1
5	-3	-3	4	3
-3	3	3	-4	3
5	2	2	3	3

feature map

+ padding  $p_h$  and  $p_w$

# Padding & Stride

0	0	0	0	0	0	0	0
0	0	1	1	0	0	0	0
0	0	1	1	0	0	0	0
0	1	0	0	1	1	0	0
0	0	1	1	0	1	0	0
0	1	1	1	1	1	0	0
0	0	0	0	0	0	0	0

image

\*

0	-1	0
-1	5	-1
0	-1	0

kernel

=

-1	3	3	-1	0
-2	3	3	-2	-1
5	-3	-3	4	3
-3	3	3	-4	3
5	2	2	3	3

feature map

+ padding  $p_h$  and  $p_w$

# Padding & Stride

0	0	0	0	0	0	0
0	0	1	1	0	0	0
0	0	1	1	0	0	0
0	1	0	0	1	1	0
0	0	1	1	0	1	0
0	1	1	1	1	1	0
0	0	0	0	0	0	0

image

\*

0	-1	0
-1	5	-1
0	-1	0

kernel

=

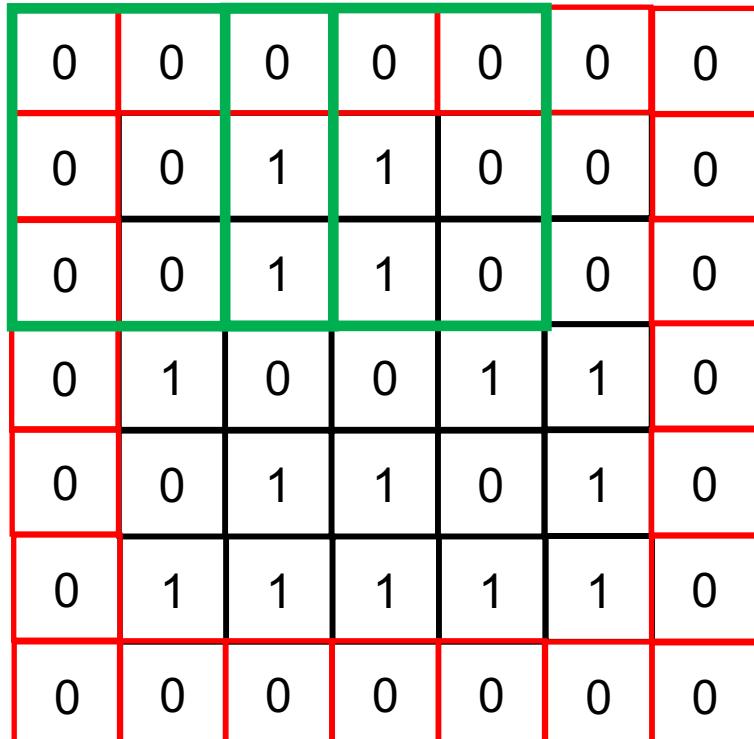
-1	3	3	-1	0
-2	3	3	-2	-1
5	-3	-3	4	3
-3	3	3	-4	3
5	2	2	3	3

feature  
map

+ padding  $p_h$  and  $p_w$

# Padding & Stride

2 steps!



0	0	0	0	0	0	0
0	0	1	1	0	0	0
0	0	1	1	0	0	0
0	1	0	0	1	1	0
0	0	1	1	0	1	0
0	1	1	1	1	1	0
0	0	0	0	0	0	0

\*

0	-1	0
-1	5	-1
0	-1	0

kernel

=

-1	3	0
5	-3	3
5	2	3

feature  
map

stride = 2

image

Linear Layer:

$$\mathbf{a}^{[l]} = g^{[l]} \left( (\mathbf{W}^{[l]})^T \mathbf{a}^{[l-1]} \right)$$

Convolution  
Layer:

$$\mathbf{A}^{[l]} = g^{[l]} \left( \mathbf{W}^{[l]} * \mathbf{A}^{[l-1]} \right)$$

Weights = **Kernels**  
(3D matrix)

could be  
same

2D matrix

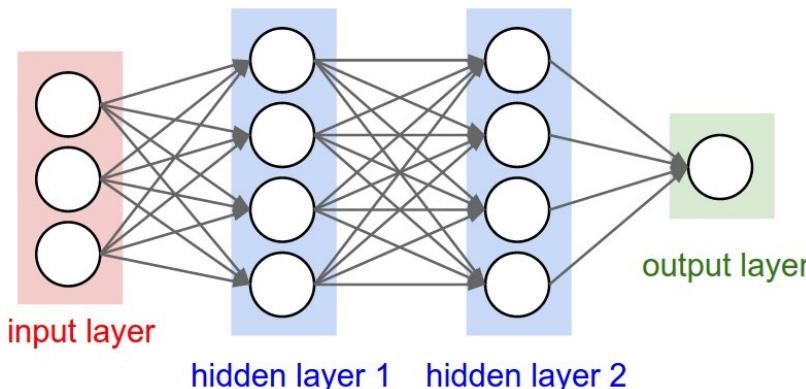
/

**Activation** is a 3D Matrix!  
Concatenation of Feature  
Maps

# MLP vs CNN

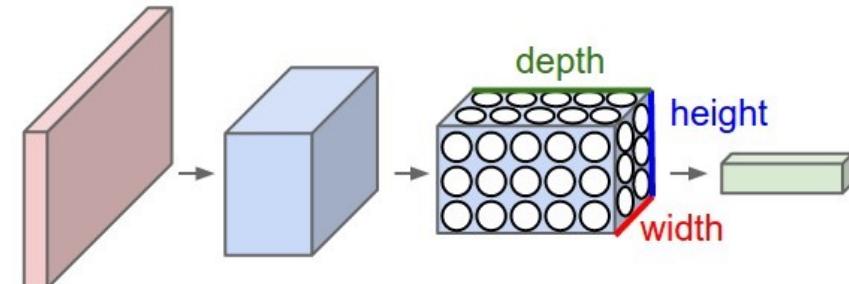
## Fully Connected Linear Layers

- Each layer has multiple **neurons**
- Neuron output: **0D scalar activation**
- Neuron input: **1D vector of activations**
  - Each *element* is a different neuron
- Each layer is a **1D vector**



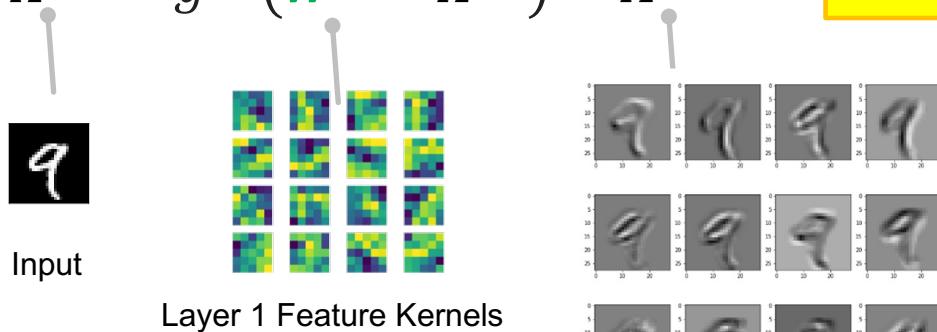
## Convolutional Layers

- Each layer has multiple **ernels**
- Kernel output: **2D matrix feature map**
- Kernel input: **3D matrix of feature maps**
  - Each *depth position* is a different kernel
  - Analogy: filters are “stacked” together
- Each layer is a **3D matrix**



# Convolution Layer

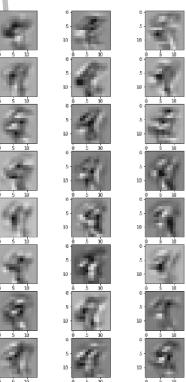
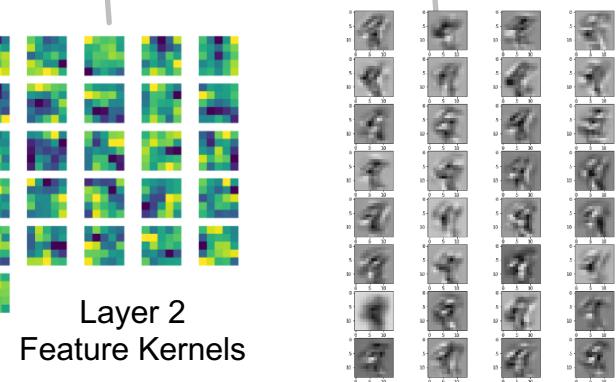
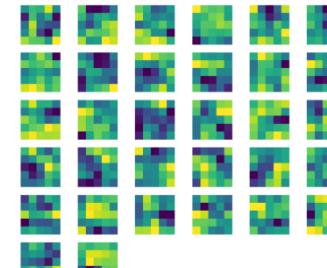
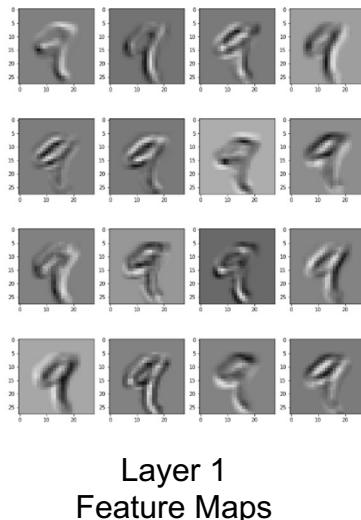
$$X^{[0]} \rightarrow g^{[1]}(W^{[1]} * X^{[0]}) = A^{[1]} \xrightarrow{\text{Pooling}} g^{[2]}(W^{[2]} * X^{[1]}) = A^{[2]}$$



## Hyperparameters

1. Number of kernels  $k$
2. Kernel size  $\kappa$
3. Padding  $p$
4. Stride  $s$

Chosen manually, or automatically with [hyperparameter tuning](#)



Kernels are learned *automatically* through **weight updates**.

**Interpretability:** do you know what these **kernel** mean?

# Feature Detectors: Intuition of Neuron Kernels in Layers

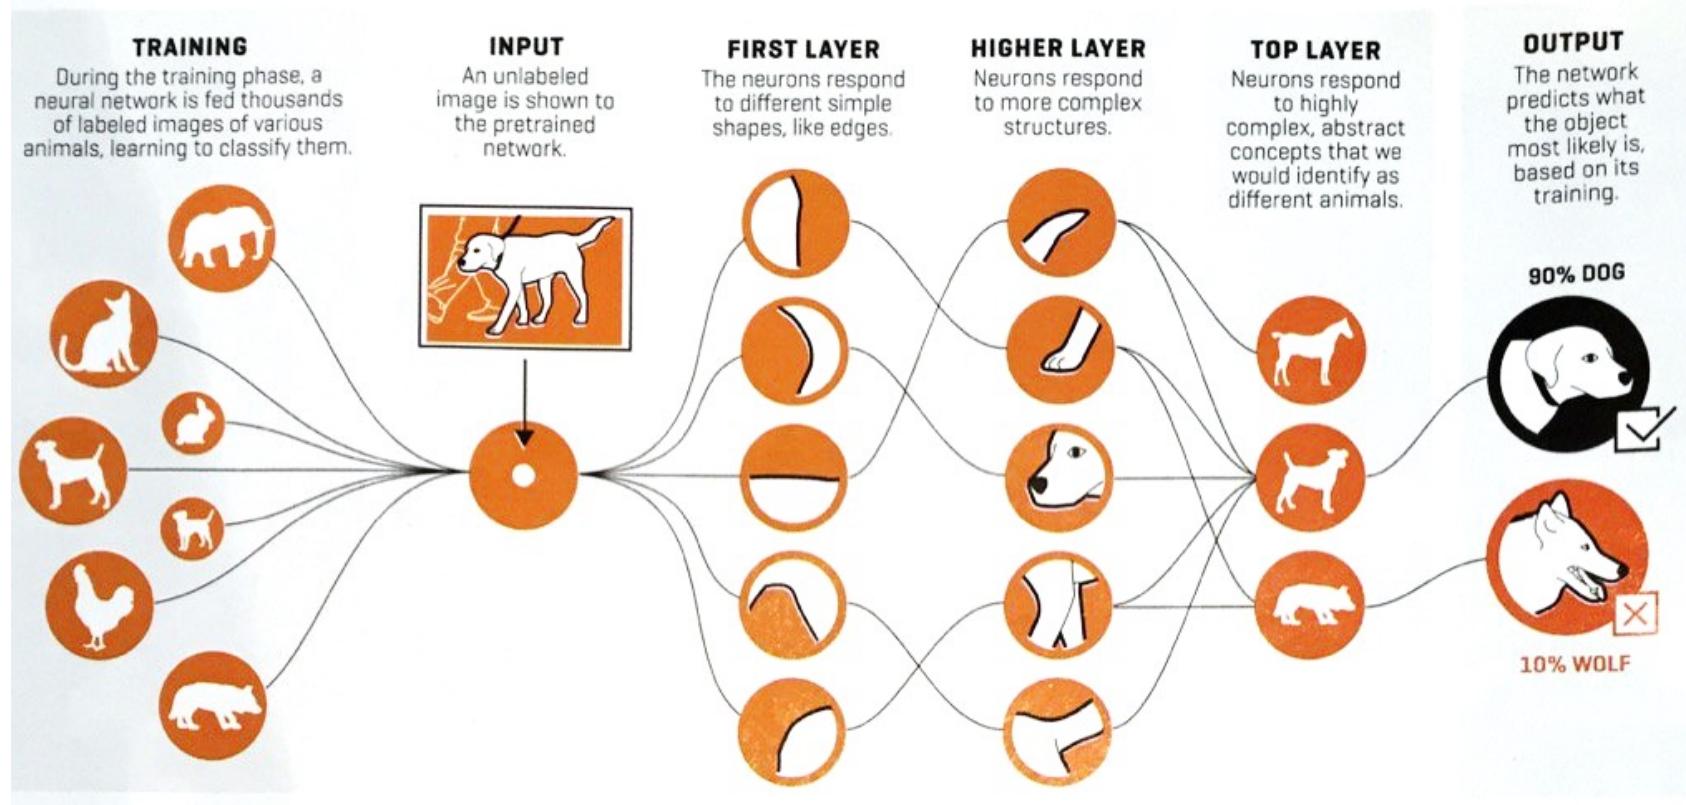


Image credit: <https://fortune.com/longform/ai-artificial-intelligence-deep-machine-learning/>

**Analogy:** activations of different filters learned by CNNs is like seeing the image through different lens filters



Image credit: <https://www.amazon.com/Godefa-Samsung-Andriod-Smartphone-Universal/dp/B07RQRLQYH>

<https://www.yankodesign.com/2020/02/17/this-retro-inspired-camera-records-dreamy-looking-gifs-that-replicate-vintage-8mm-film/>

# Convolutional Neural Network

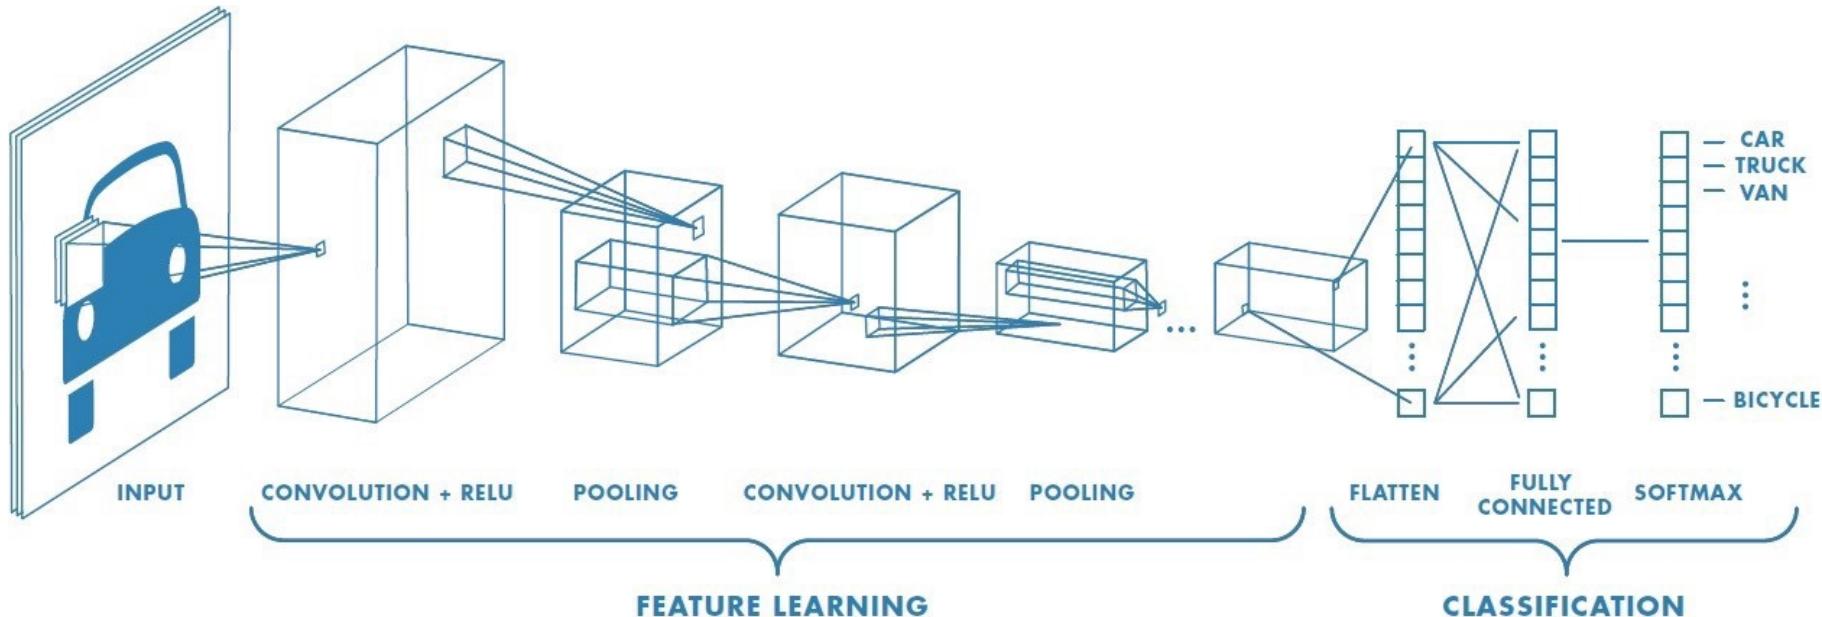
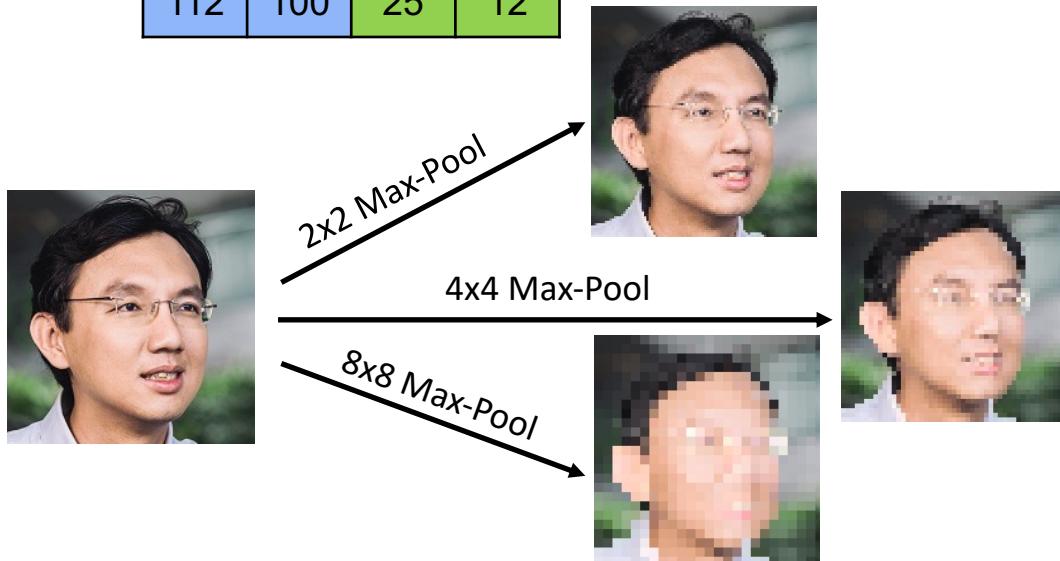
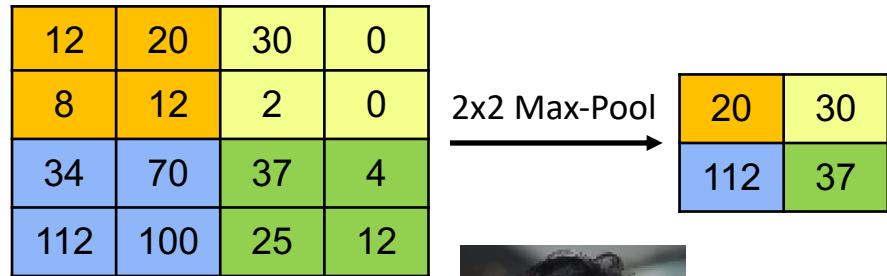


Image credit: <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

# Pooling Layer

- **Downsamples Feature Maps**
- Helps to train later ernels to detect **higher-level** features
- Reduces **dimensionality**
- Aggregation methods
  - Max-Pool (most common)
  - Average-Pool
  - Sum-Pool



# How do we decide on the final answer?

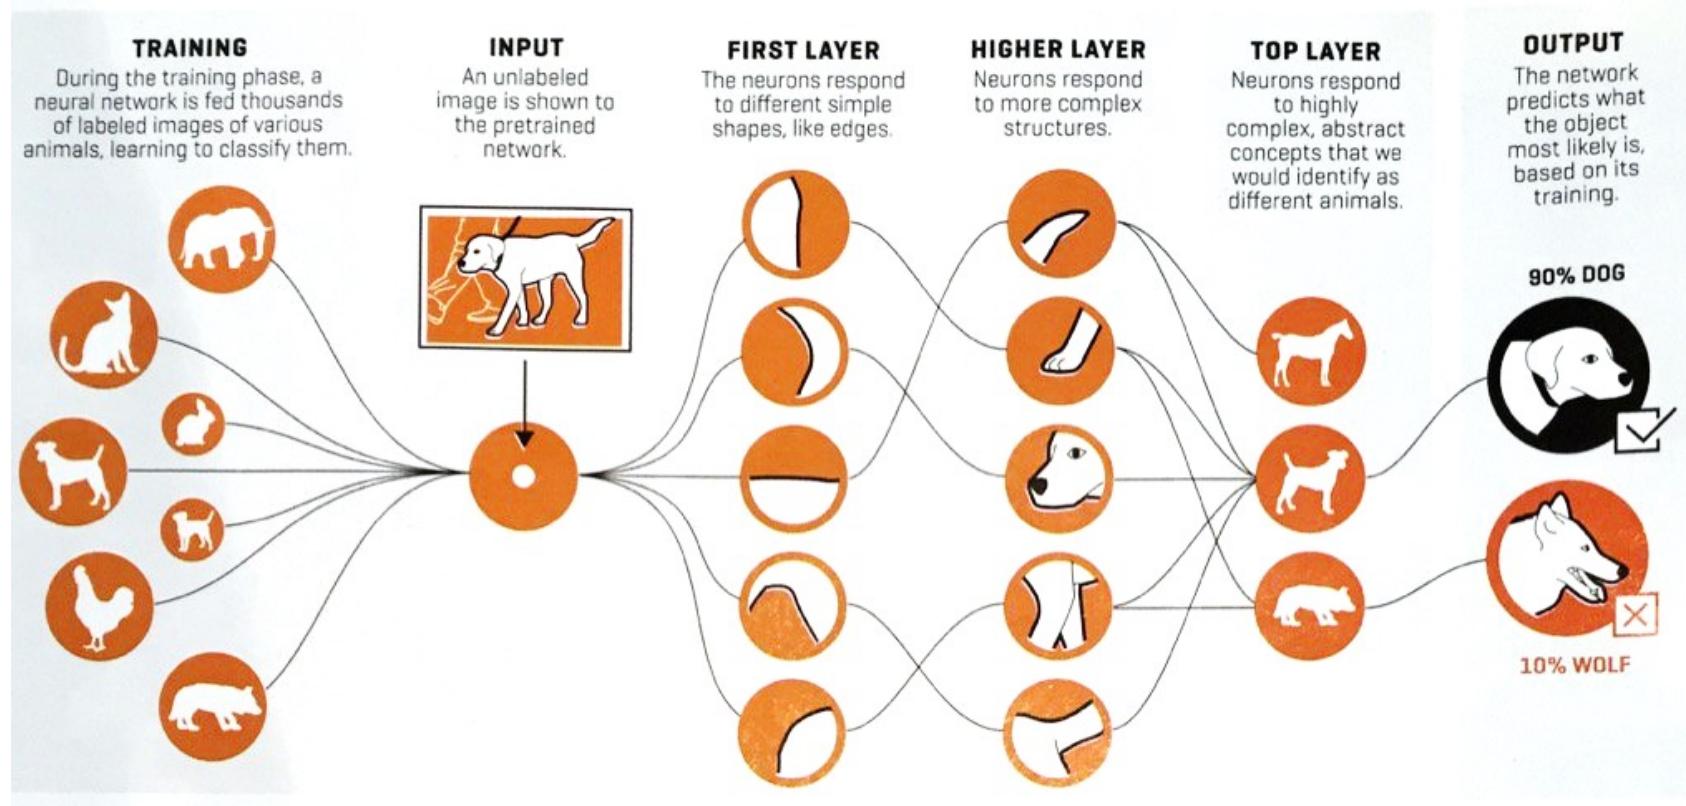


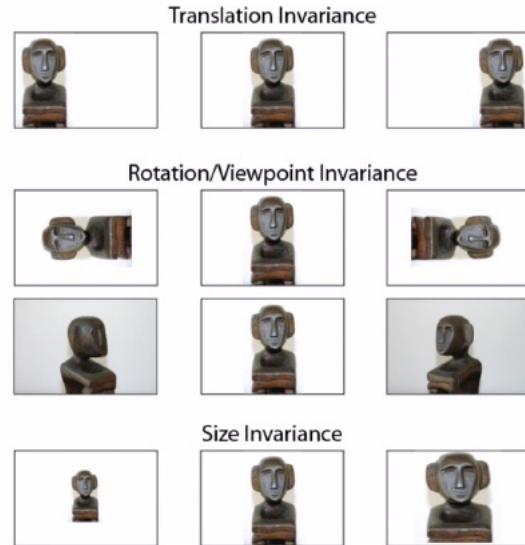
Image credit: <https://fortune.com/longform/ai-artificial-intelligence-deep-machine-learning/>

# Why is convolution better?

- Sparse Connections:
  - Each output is connected only to inputs within receptive field vs all inputs
  - Fewer parameters
  - Less overfitting
- Weight sharing vs unique weights
  - Regularization
  - Less overfitting
- Location or Spatial Invariance
  - Function transformation should not depend on location within the image
  - Make the same prediction no matter where the object is in the image

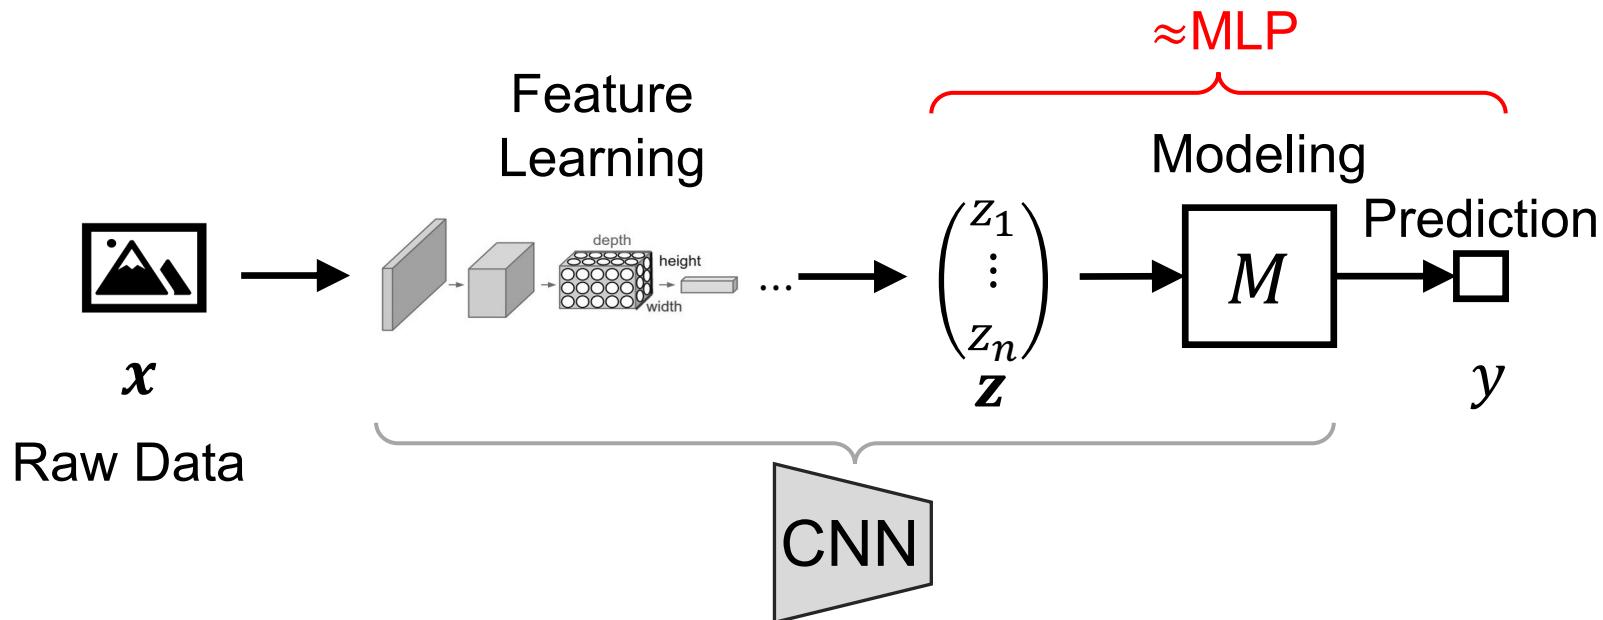
# Location or Spatial Invariant

- You can recognize an object even if its appearance varies in some way
- Convolution operation commutes with respect to translation
  - If you convolve  $f$  with  $g$ , it doesn't matter if you translate the convolved output  $f * g$ , or you translate  $f$  or  $g$  first, then convolve them.
  - <https://en.wikipedia.org/wiki/Convolution>

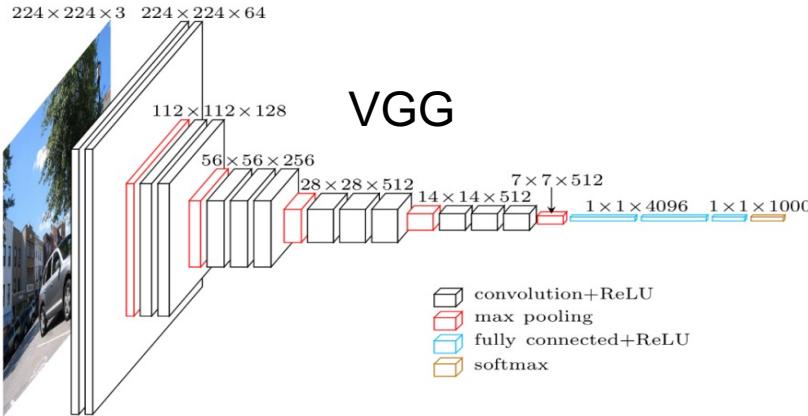


Slide credit: Matt Krause

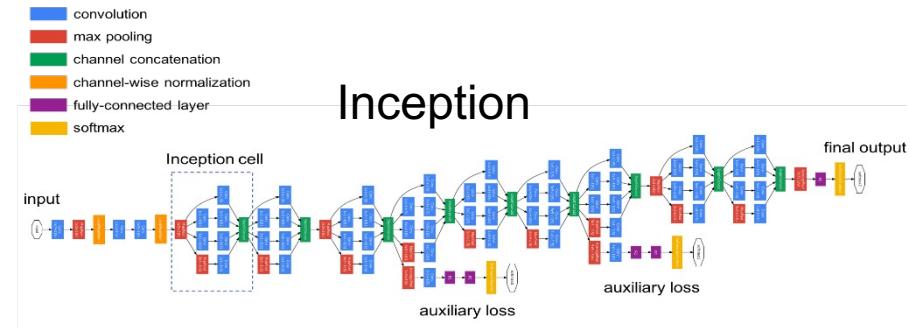
# Summary: CNN



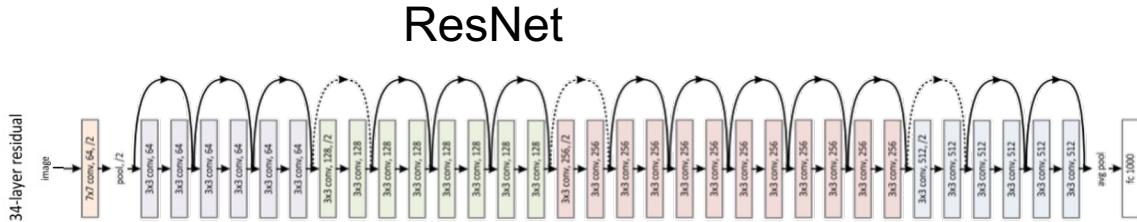
# Other popular CNN architectures



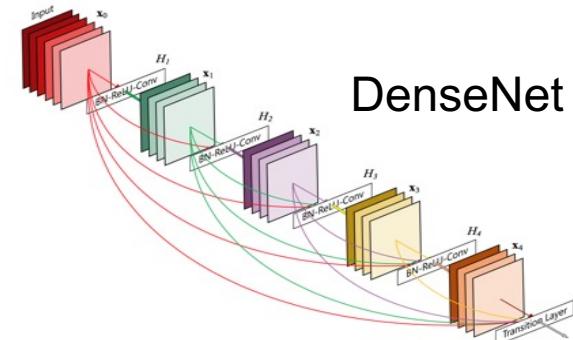
VGG



Inception



ResNet



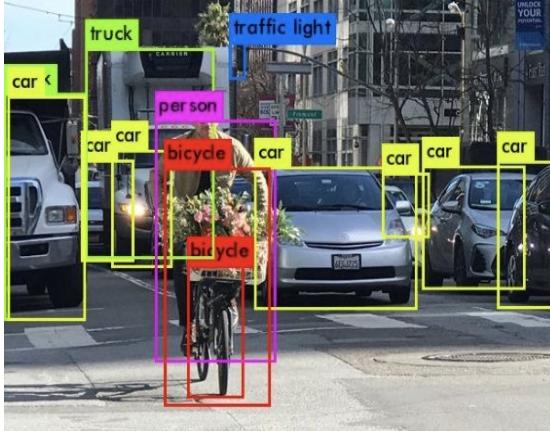
DenseNet

Further reading: <https://www.jeremyjordan.me/convnet-architectures/>

# Applications of CNN



Image Classification  
e.g., face emotions



Object Detection  
e.g., self-driving cars

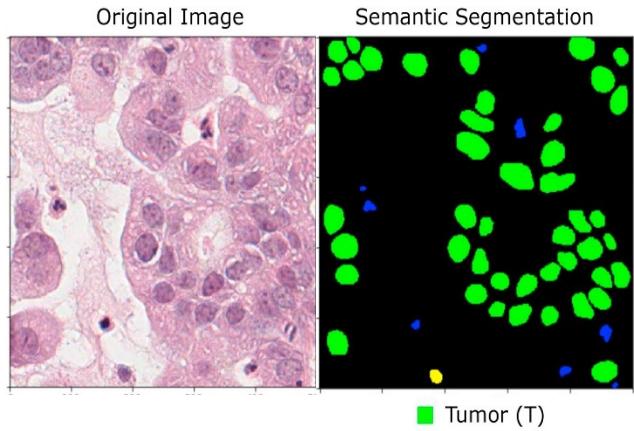


Image Segmentation  
e.g., cancer cell detection

Image credit:

<https://monica-dommaraju.medium.com/analysis-of-deep-learning-based-object-detection-f14d5138148>

[https://ajp.amjpathol.org/article/S0002-9440\(18\)31121-0/fulltext](https://ajp.amjpathol.org/article/S0002-9440(18)31121-0/fulltext)

<https://appliedmachinelearning.blog/2018/11/28/demonstration-of-facial-emotion-recognition-on-real-time-video-using-cnn-python-keras/>