




# IT5005 Artificial Intelligence

Sirigina Rajendra Prasad  
AY2025/2026: Semester 1

Markov Decision Processes (MDP)

# Greedy best-first search - without heuristics

- Cost of each step is 1



# Greedy best-first search - without heuristics

- Initialize heuristic cost with zero for all states
  - $h(s) = 0 \forall s$
- Assume neighbors have perfect information on cost of reaching the goal at each iteration
- Repeat the following until there is no change in  $h(s)$ :
  - $h(s) = \min(h(s') + c(s, s')) \forall s$ , where  $s' \in N(s)$ 
    - where,  $N(s)$  is the set of neighbours of state  $s$  and  $c(s, s')$  is the cost of moving from state  $s$  to state  $s'$
- Select the action at each state through greedy best-first search

# Greedy best-first search - without heuristics

- Initialize heuristic cost with zero
- Assume neighbors have perfect information on cost of reaching the goal
- Update the cost of reaching the goal from each state iteratively
  - At each iteration select the neighbor with least cost

$$h_k(s) = \min(h_{k-1}(s') + c(s, s')) \quad \forall s, \text{ where } s' \in N(s) \text{ and } c(s, s') = 1$$

Initialization:

0 ↓ 0	0 ↑ 0	0 ↓ 0
0 ↑ 0	0 ↓ 0	0 ↑ 0
Start 0	Goal 0	

$$h_0(s) = 0 \quad \forall s$$

Iteration #1

1 ↓ 1	1 ↑ 1	1 ↓ 1
1 ↑ 1	1 ↓ 1	1 ↓
Start 1	Goal 0	

Iteration #2

2 ↓ 2	2 ↓ 2	2 ↓
2 ↑ 2	2 → 2	1 ↓
Start 2	Goal 0	

Iteration #3

3 ↓ 3	3 ↓ 2	2 ↓
3 ↑ 2	2 → 1	1 ↓
Start 3	Goal 0	

Iteration #4

4 ↓ 3	3 ↓ 2	2 ↓
3 → 2	2 → 1	1 ↓
Start 4	Goal 0	


Iteration #5

4 ↓ 3	3 ↓ 2	2 ↓
3 → 2	2 → 1	1 ↓
Start 4	Goal 0	


No change

# Greedy Search

without uncertainty



with uncertainty



Select a neighbor with minimum heuristic cost

Select a neighbor with minimum **average** heuristic cost

**“expectimin”**

$P(s_j|s_i, a_k)$ : Probability that action  $a_k$  at state  $s_i$  leads to state  $s_j$


# Dynamic Programming (DP)

- DP can be used if a problem can be divided into subproblems and the problem has either **overlapping subproblems** or **optimal substructure**
- Overlapping subproblems
  - Stores solutions of subproblems to avoid repeated calculations
    - Ex: Fibonacci sequence
- Optimal substructure
  - Obtains optimal solution of a problem by using optimal solutions of its subproblems

# Why didn't we consider DP for search?

Not suitable for state spaces with large number of states


State Space for Uninformed and Informed Search




# So far: (1) Reasoning in State Space

- Based on Atomic Representation
- Uninformed and Informed Search
- a sequence of **actions** that take the agent from initial state to goal state with minimal **cost**

State Space for Local Search




State Space for Adversarial Search



# So far: (2) Reasoning in Feature Space

- Represent each state using features/variables
  - Discrete and Boolean variables
- Reason using features
  - Efficient search through
    - Consistency checking
    - Heuristics
  - Automated answering of queries using
    - Resolution Rule and Contradiction
    - Modus Ponens



# So far: (3) Reasoning in Uncertain Environment

- Efficient Representation of knowledge using Bayesian Networks, Markov Chains, and Hidden Markov Models
- Identification of relevant variables via independence relations
- Answering queries using exact Inference

# Markov Decision Process

Uncertainty in Action Effects

# Markov Decision Process

“A sequential decision problem for a fully observable, stochastic environment with a Markovian transition model and additive rewards”

# MDP Applications

- Self-driving car
  - Taking actions (steering, accelerating, breaking, etc.) to **optimize** performance (eg: ride quality)
- Games
  - Movement of pieces to **maximize** chances of winning
- Logistics planning
  - Inventory movements to **minimize** delays in delivery
- Humanoid robot
  - Movement planning to **optimize** time to complete a task
- Managing investment portfolio
  - Trades to **optimize** long-term investment gains
- Wireless networks
  - Scheduling and resource allocation for **optimal utilization** of resources

# MDP

- Isn't it an optimization framework?
  - Yes, provided the MDP model is known
- Then, why are we doing it in AI Module?
  - Foundation for reinforcement learning (RL)
  - If MDP model is unknown, agent learns the optimal actions by acting and collecting the rewards in the environment

# Markov Decision Process


- An MDP is defined as:
  - A set of states
  - A set of actions  $ACTIONS(s)$  for each state  $s$
  - Transition Function -  $P(s'|s, a)$
  - Reward Function –  $R(s, a, s')$
  - Discount Factor –  $\gamma$

$P(s'|s, a)$ : the probability of going from state  $s$  to state  $s'$  with action  $a$   
 $R(s, a, s')$ : Reward for the transition from state  $s$  to state  $s'$  with action  $a$

For each state  $s$  and action  $a$ :


$$\sum_{s'} P(s'|s, a) = 1$$

# Self-driving Car MDP




# Traditional Models Vs MDPs

Traditional Model (no uncertainty)



MDP (Effect Uncertainty)




Traditional Models can be treated as special case of MDPs

$$P(s_2|s_1, a_1) = P(s_3|s_1, a_2) = 1$$

$$R(s_1, a_1, s_2) = -c_1$$

$$R(s_1, a_1, s_3) = -c_2$$

# Maze - Traditional Model



## State Space Modeling:

**Initial State:** (1,1)

**Actions:**  $E, W, N, S$

**Transition Model:** Example: Action  $E$  takes agent from (1,1) to (2,1)


**Step-Cost:** 1

**Goal-Test:** State (4,3)



Search algorithms can find  
optimal solution (path with least cost)


# Maze - MDP Model



$$\begin{aligned} P((1,1)|(1,1), \rightarrow) &= 0.1 \\ P((2,1)|(1,1), \rightarrow) &= 0.8 \\ P((1,2)|(1,1), \rightarrow) &= 0.1 \\ P((1,3)|(1,1), \rightarrow) &= 0 \end{aligned}$$

$$P(s'|s, a)$$

(Transition Function)



$$\begin{aligned} R((1,1), \rightarrow, (1,1)) &= -0.04 \\ R((1,1), \rightarrow, (2,1)) &= -0.04 \\ R((1,1), \rightarrow, (1,2)) &= -0.04 \\ R((3,3), \rightarrow, (4,3)) &= +1 \\ R((3,2), \rightarrow, (4,2)) &= -1 \\ R((4,1), \rightarrow, (4,2)) &= -1 \end{aligned}$$

$$\begin{aligned} R(s, a, s') \\ (\text{Reward Function}) \end{aligned}$$

## Effect Uncertainty:

Desired effect with probability 0.8  
Undesired effects with probability 0.2

Transition to State	Reward
Terminal State (4,3)	+1
Terminal State (4,2)	-1
Other states	-0.04

## Example:


Action  $\uparrow$  at (1,1):

Move to (1,2) with probability 0.8  
Stay at (1,1) with probability 0.1  
Move to (2,1) with probability 0.1

Action  $\leftarrow$  at (1,1):

Move to (1,2) with probability 0.1  
Stay at (1,1) with probability 0.9

# Traditional Model - Search Tree





Goal-based agent:

Search for a path to reach goal state (assume (4,3) is the goal state)

# MDP Search Tree




- Move East
- ← Move West
- ↑ Move North
- ↓ Move South



# A Closer Look at MDP Search Tree

- For state (1,1) and action  $\rightarrow$ :



$$P((1,1)|(1,1), \rightarrow) = 0.1$$

$$P((2,1)|(1,1), \rightarrow) = 0.8$$


$$P((1,2)|(1,1), \rightarrow) = 0.1$$

$$R((1,1), \rightarrow, (1,1)) = -0.04$$

$$R((1,1), \rightarrow, (2,1)) = -0.04$$

$$R((1,1), \rightarrow, (1,2)) = -0.04$$

$$\sum_{s'=\{(1,1),(2,1),(1,2)\}} P(s'|1,1, \rightarrow) = 1$$



# MDP Terminology

- Policies
  - Policy ( $\pi$ )
  - Optimal Policy ( $\pi^*$ )
- Utilities
  - $U^\pi(s)$ : Utility of a state  $s$  with policy  $\pi$
  - $U(s)$ : Optimal utility of a State
- Action-Utility Function or Q-function
  - $Q^\pi(s, a)$ : Action-Utility value by taking action  $a$  at state  $s$  and then following policy  $\pi$
  - $Q(s, a)$ : Action-Utility value by taking action  $a$  at state  $s$  and then following optimal policy

# Policy


- Policy  $\pi(s)$  suggests an action  $a$  for each state  $s$

- $\pi: s \rightarrow a$

$$\begin{aligned}\pi_1((1,1)) &= \uparrow \\ \pi_1((2,1)) &= \rightarrow \\ \pi_1((1,2)) &= \uparrow\end{aligned}$$


- Total number of policies:  $|A|^{|S|}$ 
  - $|A|$  is the number of actions
  - $|S|$  is the number of states

An Example:  $\pi_1$




# Policy Examples


Policy 1 ( $\pi_1$ )



Policy 2 ( $\pi_2$ )




Policy 3 ( $\pi_3$ )



## Questions:

- Which policy is good (optimal)?
- How to define optimality?
- How to find optimal policy?



# Utility of a State with a Policy $\pi$ : $U^\pi(s)$

Expected (average) cumulative sum of rewards obtained by following the policy  $\pi$  from that state

Utility of a state is also called State-Value or Value of a State

# Cumulative Sum of Rewards

- Sequence of states and actions with policy  $\pi$  from the initial state  $s_0$

$$[s_0, \pi(s_0), s_1, \pi(s_1), \dots, \pi(s_{n-1}), s_n, \dots]$$



Trajectory with policy  $\pi$  for an episode


$$\begin{aligned}\pi: s \rightarrow a \\ \pi(s_i) = a_i\end{aligned}$$

- Sequence ends at terminal state
  - Sequence is also known as trajectory of an episode
- 
- Cumulative sum of rewards for an episode with policy  $\pi$  from the initial state  $s_0$ :

$$U_h([s_0, \pi(s_0), s_1, \pi(s_1), \dots, \pi(s_{n-1}), s_n, \dots]) = R(s_0, \pi(s_0), s_1) + R(s_1, \pi(s_1), s_2) + R(s_2, \pi(s_2), s_3) + \dots$$



$$= \sum_{t=0}^{\infty} R(S_t, \pi(S_t), S_{t+1})$$



Utility with Policy  $\pi$

reward at state  $S_t$  with action  $\pi(S_t)$  and transition to state  $S_{t+1}$

# Cumulative Sum of Discounted Rewards

- Cumulative sum of discounted rewards for an episode:

$$U_h([s_0, \pi(s_0), s_1, \pi(s_1), \dots, \pi(s_{n-1}), s_n \dots]) = R(s_0, \pi(s_0), s_1) + \gamma R(s_1, \pi(s_1), s_2) + \gamma^2 R(s_2, \pi(s_2), s_3) + \dots$$

$$= \sum_{t=0}^{\infty} \gamma^t R(S_t, \pi(S_t), S_{t+1})$$

$0 \leq \gamma \leq 1$




(discount factor)

- Why discounts?

- Guarantees convergence of iterative algorithms
- Also guarantee convergence for infinite/indefinite horizon problems
- Gives more weightage to immediate reward at state  $s_0$

# Example: Maze - Cumulative Rewards



- Trajectories are **stochastic** and **not unique** for a given policy
  - Different trajectories are possible for a given policy
- **Cumulative reward** is also **stochastic**
  - Different for each episode



# Environmental History of State $(1,1)$ with Policy $\pi_1$

- States and Actions after executing the policy  $\pi_1$ 
  - $[s_0, \pi_1(s_0), s_1, \pi_1(s_1), \dots, \pi_1(s_{n-1}), s_n, \dots]$

- Example:
  - Environmental history of four episodes for  $s_0 = (1,1)$ 
    - ❶  $[(1,1), \uparrow, (1,2), \uparrow, (1,3), \rightarrow, (2,3), \rightarrow, (3,3), \rightarrow, (4,3)]$
    - ❷  $[(1,1), \uparrow, (1,1), \uparrow, (2,1), \rightarrow, (3,1), \uparrow, (3,2), \uparrow, (3,3) \rightarrow, (4,3)]$
    - ❸  $[(1,1), \uparrow, (1,1), \uparrow, (1,2), \uparrow, (1,3), \rightarrow, (2,3), \rightarrow, (3,3), \rightarrow, (4,3)]$
    - ❹  $[(1,1), \uparrow, (1,1), \uparrow, (2,1), \rightarrow, (3,1), \uparrow, (3,2), \uparrow, (4,2)]$
  - Each episode has a different trajectory with same policy  $\pi_1$



# Cumulative Rewards of State (1,1) with Policy $\pi_1$

- States and Actions after executing the policy  $\pi_1$

$$[s_0, \pi_1(s_0), s_1, \pi_1(s_1), \dots, \pi_1(s_{n-1}), s_n, \dots]$$

- Examples:

- Episode 1

- Trajectory:  $[(1,1), \uparrow, (1,2), \uparrow, (1,3), \rightarrow, (2,3), \rightarrow, (3,3), \rightarrow, (4,3)]$
- Reward:  $-0.04 - \gamma 0.04 - \gamma^2 0.04 - \gamma^3 0.04 + \gamma^4 1$

- Episode 2

- Trajectory:  $[(1,1), \uparrow, (1,1), \uparrow, (2,1), \rightarrow, (3,1), \uparrow, (3,2), \uparrow, (3,3), \rightarrow, (4,3)]$
- Reward:  $-0.04 - \gamma 0.04 - \gamma^2 0.04 - \gamma^3 0.04 - \gamma^4 0.04 + \gamma^5 1$



- Episode 3

- Trajectory:  $[(1,1), \uparrow, (1,1), \uparrow, (1,2), \uparrow, (1,3), \rightarrow, (2,3), \rightarrow, (3,3), \rightarrow, (4,3)]$
- Reward:  $-0.04 - \gamma 0.04 - \gamma^2 0.04 - \gamma^3 0.04 - \gamma^4 0.04 + \gamma^5 1$

- Cumulative rewards are random

- Questions:

- What is the **average cumulative reward** of state (1,1) with policy  $\pi_1$ ?
- What are the **probabilities** of these rewards?



# Cumulative Sum of Discounted Rewards

- A recursive representation of cumulative sum of discounted rewards

$$U_h([s_0, \pi(s_0), s_1, \pi(s_0), \dots, \pi(s_{n-1}), s_n \dots]) = R(s_0, \pi(s_0), s_1) + \gamma [R(s_1, \pi(s_1), s_2) + \gamma R(s_2, \pi(s_2), s_3) + \dots]$$



Reward of a specific episode with initial state  $s_0$  and policy  $\pi$



It is a random value and changes with episode



We need average cumulative reward

$$= R(s_0, \pi(s_0), s_1) + \gamma U_h([s_1, \pi(s_1), \dots, s_n, \dots])$$

Rewards due to transition from current state ( $s_0$ ) to next state ( $s_1$ ) with action  $\pi(s_0)$

Cumulative sum of rewards from next state ( $s_1$ ) onwards if the policy  $\pi$  is followed

# Utility of State $s_0$ with a Policy $\pi$

- Expected (average) utility of a state  $s_0$  with policy  $\pi$


$$U^\pi(s_0) = E \left[ \sum_{t=0}^{\infty} \gamma^t R(S_t, \pi(S_t), S_{t+1}) \right] \quad 0 \leq \gamma \leq 1$$

- Expectation is with respect to probability distribution of state sequences (determined by  $s_0$  and  $\pi$ )
  - Infinite dimensional expectation
  - Markov assumption simplifies the above expectation

# Markov Condition

- Given the current state and action, next state is independent of previous states and actions


$$\Pr(s_{t+1}|s_t, a_t, s_{t-1}, a_{t-1}, \dots) = \Pr(s_{t+1}|s_t, a_t)$$



# Markov Condition Enables....

- Factoring of joint distribution

$$\Pr(s_0, a_0, s_1, a_1, s_2, a_2, s_3) = \Pr(s_3|s_2, a_2) \Pr(s_2|s_1, a_1) \Pr(s_1|s_0, a_0) \Pr(s_0)$$



- Decomposition of complex problem into small sub-problems

Deterministic Policy:  
 $\Pr(a_i|s_i) = 1$

# Expected Utility of State $s_0$ with a Policy $\pi$ : $U^\pi(s_0)$

$$U^\pi(s_0) = E \left[ \sum_{t=0}^{\infty} \gamma^t R(S_t, \pi(S_t), S_{t+1}) \right]$$



Only if Markov conditions are satisfied

$$U^\pi(s_0) = \sum_{s'} P(s'|\pi(s_0), s_0) [R(s_0, \pi(s_0), s') + \gamma U^\pi(s')]$$

# Bellman Update Equation for Utility

- Expected utility of state  $s$  with a policy  $\pi$


$$U^\pi(s) = \sum_{s'} P(s'|\pi(s), s)[R(s, \pi(s), s') + \gamma U^\pi(s')]$$

Utility of state  $s$  with policy  $\pi$

Probability of moving to state  $s'$  with action  $\pi(s)$  at state  $s$

Reward for taking an action  $\pi(s)$  at state  $s$  and moving to state  $s'$

Utility of state  $s'$  with policy  $\pi$



$s'$  is this set of states


Utility of a state is expressed in terms of utility of neighbors

# Action-Utility or $Q$ -Value

- $Q$  value indicates expected value of an action at a state
- $Q^\pi(s, a)$ : expected utility of state  $s$  with an action  $a$  at state  $s$ , and then subsequently following the policy  $\pi$

$$Q^\pi(s, a) = \sum_{s'} P(s'|a, s)[R(s, a, s') + \gamma U^\pi(s')]$$

where  $a = \pi(s)$




# Relation between $Q^\pi(s, a)$ and $U^\pi(s)$

- If the policy  $\pi$  is fixed, utility is same as action-utility

$$U^\pi(s) = Q^\pi(s, a)$$

$$U^\pi(s) = \sum_{s'} P(s'|a, s)[R(s, a, s') + \gamma U^\pi(s')]$$

$$Q^\pi(s, a) = \sum_{s'} P(s'|a, s_1)[R(s, a, s') + \gamma U^\pi(s')]$$



Deterministic policy → Action fixed for each state → Utility of state is same as utility of state-action

# Bellman Update Equation for Action-Utility

Q-value in terms of Q-values of neighbors

$$Q^\pi(s, a) = \sum_{s'} P(s'|a, s)[R(s, a, s') + \gamma Q^\pi(s', a')]$$


$$U^\pi(s') = Q^\pi(s', a')$$

$$a = \pi(s)$$

$$a' = \pi(s')$$

$$s' = \{s_1, s_2, s_3\}$$

$$a' = \{a_1, a_2, a_3\}$$




# $U$ -value vs $Q$ -value

- $U(s)$  indicates utility or value or desirability of a state  $s$
- $Q(s, a)$  indicates value of action  $a$  at state  $s$ 
  - Can be used to select **optimal action** at state  $s$ 
    - Find value of all actions at a state
    - Select action with largest  $Q$ -value
  - In other words,  $Q$ -value is used to make **optimal decision**
- Q-learning is one of the paradigms in reinforcement learning (RL)
  - Q-learning involves finding  $Q$ -values when model (MDP) parameters are unknown
  - $Q$ -values implicitly stores optimal policy, i.e., optimal  $Q$ -values reveal optimal policy



# Maze: Expected Utility of state (1,1) with policy $\pi_1$

Policy:  $\pi_1$

Policy  $\pi_1$



$$U^{\pi_1}((1,1)) = 0.1[-0.04 + \gamma U^{\pi_1}((1,1))] + 0.8[-0.04 + \gamma U^{\pi_1}((1,2))] + 0.1[-0.04 + \gamma U^{\pi_1}((2,1))]$$




$$U^\pi(s) = \sum_{s'} P(s'|\pi(s), s)[R(s, \pi(s), s') + \gamma U^\pi(s')]$$

# Maze: Bellman Update Equations

- We can get similar equations for other states:
- Example:

$$U^{\pi_1}((1,2)) = 0.2[-0.04 + \gamma U^{\pi_1}((1,2))] + 0.8[-0.04 + \gamma U^{\pi_1}((1,3))]$$




Utilities of terminal states\*:

$$\begin{aligned} U^{\pi_1}((4,3)) &= 0 \\ U^{\pi_1}((4,2)) &= 0 \end{aligned}$$

$$\begin{aligned} \forall a P((4,2)|(4,2), a) &= 1, R((4,2), a, (4,2)) = 0 \\ \forall a P((4,3)|(4,3), a) &= 1, R((4,3), a, (4,3)) = 0 \end{aligned}$$

Policy:  $\pi_1$



\*Terminal states can also move to exit states with a specific reward:


$$\left. \begin{aligned} \forall a P(EXIT|(4,2), a) &= 1, R((4,2), a, EXIT) = -1 \\ \forall a P(EXIT|(4,3), a) &= 1, R((4,3), a, EXIT) = 1 \end{aligned} \right\} \begin{aligned} U^{\pi_1}((4,2)) &= -1 \\ U^{\pi_1}((4,3)) &= 1 \end{aligned}$$

# Inferencing with MDP

- Prediction Problem
  - Finding utility of a state given a policy  $\pi$
  - i.e., calculation of  $U^\pi(s)$
- Control Problem
  - Finding an optimal policy  $\pi^*$
  - i.e., finding an action for each state ( $s$ ) that maximizes its utility ( $U(s)$ )


$$\pi^*(s) = \arg \max_{\pi} U^\pi(s)$$

## Prediction or Policy Evaluation



Given a policy, what is the reward at each state?

# Control



What is the optimal policy?

Finding optimal action at each state such that the average reward is maximized

# Prediction Problem

- Calculation of expected utility of state  $s$  with a policy  $\pi$

Known quantities (Model Parameters)

$$U^\pi(s) = \sum_{s'} P(s'|\pi(s), s) [R(s, \pi(s), s') + \gamma U^\pi(s')]$$

unknowns

- $|S|$  linear equations with  $|S|$  unknowns
- How to solve these equations?
- Two methods:
  - Direct method using linear algebra
  - Iterative Methods

# Prediction Problem

- Calculation of Action-Utility for a given policy


Known quantities (Model Parameters)

$$Q^\pi(s, a) = \sum_{s'} P(s'|a, s_1) [R(s, a, s') + \gamma Q^\pi(s', a')]$$

unknowns

$$a = \pi(s)$$

$$a' = \pi(s')$$



$$s' = \{s_1, s_2, s_3\}$$

$$a' = \{a_1, a_2, a_3\}$$

# Direct Method: Example

Linear Equations

$$a_{11}x_1 + a_{12}x_2 + \cdots + a_{1n}x_n = b_1$$

$$a_{21}x_1 + a_{22}x_2 + \cdots + a_{2n}x_n = b_2$$

..

$$a_{n1}x_1 + a_{n2}x_2 + \cdots + a_{nn}x_n = b_n$$

$$A \mathbf{x} = \mathbf{b} \Rightarrow \mathbf{x} = A^{-1}\mathbf{b}$$



$$A = \begin{bmatrix} a_{11} & \cdots & a_{1n} \\ \vdots & \ddots & \vdots \\ a_{n1} & \cdots & a_{nn} \end{bmatrix}$$

**Complexity:**  $O(|S|^3)$

Ok for MDPs with small state spaces

Challenging for MDPs with large state spaces

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad \mathbf{b} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix}$$

# Iterative Methods

$$a_{11}x_1 + a_{12}x_2 + \dots + a_{1n}x_n = b_1$$

$$a_{21}x_1 + a_{22}x_2 + \dots + a_{2n}x_n = b_2$$

$$\vdots$$
$$a_{n1}x_1 + a_{n2}x_2 + \dots + a_{nn}x_n = b_n$$

Iterative method

$$x_{1,k} = -\frac{a_{12}}{a_{11}}x_{2,i} - \dots - \frac{a_{1n}}{a_{11}}x_{n,k-1} + \frac{b_1}{a_{11}}$$

$$x_{2,k} = -\frac{a_{21}}{a_{22}}x_{1,i} - \dots - \frac{a_{2n}}{a_{22}}x_{n,k-1} + \frac{b_2}{a_{22}}$$

$$x_{n,k} = -\frac{a_{12}}{a_{nn}}x_{2,k-1} - \dots - \frac{a_{n-1n}}{a_{nn}}x_{n-1,k-1} + \frac{b_n}{a_{nn}}$$

$$\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

- Initialize the vector  $x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$  with zeros and evaluate the above equations iteratively


- Stop when there is no change in the values or after a fixed number of iterations

- Suitable for **nonlinear** equations as well
  - We need them later in Value Iteration algorithm
- Need to take care of convergence

# So far: Prediction

- Policy Evaluation Given a Policy
- $U^\pi(s)$ : utility of state  $s$  with a given policy  $\pi$

$$U^\pi(s) = \sum_{s'} P(s'|\pi(s), s)[R(s, \pi(s), s') + \gamma U^\pi(s')]$$



# Next: Control

- Control:
  - Finding optimal policy
- How?
  - Brute Force Algorithm
  - Policy Iteration (PI)
  - Value Iteration (VI), etc.

# Finding Optimal Policy: Brute Force Algorithm

- Find utilities of states with all policies
- Policy with maximum utility at each state is optimal policy
  - $\pi^*(s) = \arg \max_{\pi} U^{\pi}(s)$
- Number of Policies:  $|A|^{|S|}$ 
  - If size of state space ( $|S|$ ) is small and number of actions per state ( $|A|$ ) are also small, we can evaluate all policies and select the policy with best utilities across all states
  - Usually  $|S|$  is large
- Evaluation of all possible policies is impossible

# Finding Optimal Policy: VI and PI

- Relies on greedy approach
- Choose the action that maximizes the reward for next step plus discounted utility of subsequent state


$$\begin{aligned}\pi^*(s) &= \arg \max_{a \in A(s)} \sum_{s'} P(s'|a, s)[R(s, a, s') + \gamma U(s')] \\ &= \arg \max_{a \in A(s)} Q(s, a)\end{aligned}$$

(one-step look-ahead)

- But we should know optimal utility of subsequent states (i.e.,  $s'$ )!
  - Otherwise, greedy approach is suboptimal
- That means we should also know the optimal action for subsequent states!!

Find optimal action from  $A(s)$  for all  $s$

$$A(s) = \{a_1, a_2, \dots, a_N\}$$



# Optimal Utility

- Utility (value) of a state  $U(s)$  is the expected reward for the next transition plus the discounted utility of next state, assuming that the agent chooses the optimal action

Bellman equation for utility

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s'|a, s)[R(s, a, s') + \gamma U(s')]$$

$$Q(s, a) = \sum_{s'} P(s'|a, s)[R(s, a, s') + \gamma U(s')]$$

$$U(s) = \max_{a \in A(s)} Q(s, a)$$

# Optimal Action-Utility


- Action-Utility  $Q(s, a)$  of a state-action pair  $(s, a)$  is the expected reward for action  $a$  at state  $s$ , assuming that the agent chooses the optimal actions at subsequent states

$$Q(s, a) = \sum_{s'} P(s'|a, s)[R(s, a, s') + \gamma U(s')]$$

Bellman equation for Q-value

$$Q(s, a) = \sum_{s'} P(s'|a, s) [R(s, a, s') + \gamma \max_{a' \in A(s')} Q(s', a')]$$

$$\pi^*(s) = \arg \max_{a \in A(s)} Q(s, a)$$



# Value Iteration (VI)


- Start with  $U_i(s) = 0$  for all  $s$
- Calculate Q values and update Utility values

$$Q_{i+1}(s, a) \leftarrow \sum_{s'} P(s'|a, s)[R(s, a, s') + \gamma U_i(s')]$$


$$U_{i+1}(s) = \max_{a \in A(s)} Q_{i+1}(s, a) \quad \text{Bellman update equations}$$

- Once utility values converge to optimal values, select the corresponding policy

$$\pi^*(s) = \arg \max_{a \in A(s)} \sum_{s'} P(s'|a, s)[R(s, a, s') + \gamma U(s')]$$



# Maze: Bellman Equation for state (1,1)



$$Q_k((1,1), \rightarrow) = 0.8(-0.04 + \gamma U_{k-1}((1,1))) + 0.1(-0.04 + \gamma U_{k-1}((2,1))) + 0.1(-0.04 + \gamma U_{k-1}((1,2)))$$

$$Q_k((1,1), \leftarrow) = 0.9(-0.04 + \gamma U_{k-1}((1,1))) + 0.1(-0.04 + \gamma U_{k-1}((1,2))),$$

$$Q_k((1,1), \uparrow) = 0.8(-0.04 + \gamma U_{k-1}((1,2))) + 0.1(-0.04 + \gamma U_{k-1}((2,1))) + 0.1(-0.04 + \gamma U_{k-1}((1,1)))$$

$$Q_k((1,1), \downarrow) = 0.9(-0.04 + \gamma U_{k-1}((1,1))) + 0.1(-0.04 + \gamma U_{k-1}((2,1))),$$

$$U_k((1,1)) = \max\{Q_k((1,1), \rightarrow), Q_k((1,1), \leftarrow), Q_k((1,1), \uparrow), Q_k((1,1), \downarrow)\}$$

# Value Iteration (VI) Algorithm

**function** VALUE-ITERATION( $mdp, \epsilon$ ) **returns** a utility function

**inputs:**  $mdp$ , an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$ , rewards  $R(s, a, s')$ , discount  $\gamma$

$\epsilon$ , the maximum error allowed in the utility of any state

**local variables:**  $U, U'$ , vectors of utilities for states in  $S$ , initially zero

$\delta$ , the maximum relative change in the utility of any state

**repeat**

$U \leftarrow U'; \delta \leftarrow 0$

**for each** state  $s$  **in**  $S$  **do**

$U'[s] \leftarrow \max_{a \in A(s)} Q\text{-VALUE}(mdp, s, a, U)$  → Bellman update equation

**if**  $|U'[s] - U[s]| > \delta$  **then**  $\delta \leftarrow |U'[s] - U[s]|$  → Smallest change in  $U$  values among all states

**until**  $\delta \leq \epsilon(1 - \gamma)/\gamma$  → Checking for convergence

**return**  $U$



**function**  $Q\text{-VALUE}(mdp, s, a, U)$  **returns** a utility value  
return  $\sum_{s'} P(s'|a, s_1)[R(s, a, s') + \gamma U(s')]$

VI algorithm calculates optimal values

Perform **one-step look-ahead** to extract optimal policy from optimal values

$$\pi^*(s) = \arg \max_{a \in A(s)} \sum_{s'} P(s'|a, s)[R(s, a, s') + \gamma U(s')]$$

# Value Iteration: Analysis




$$N = \left\lceil \frac{\log\left(\frac{2R_{max}}{\epsilon(1-\gamma)}\right)}{\log\left(\frac{1}{\gamma}\right)} \right\rceil$$


Converges faster for small discount factor

# Optimal Policies: Example

Depends on Reward and Transition Probability




(a)



(b)

**Figure 17.2** (a) The optimal policies for the stochastic environment with  $r = -0.04$  for transitions between nonterminal states. There are two policies because in state (3,1) both *Left* and *Up* are optimal. (b) Optimal policies for four different ranges of  $r$ .

# Value Iteration: Analysis



$$\text{Policy Loss} = ||U^{\pi_i} - U||$$

$$\text{Max Error} = ||U_i - U||$$

$$\|U\| = \max_s |U(s)|$$

$U^{\pi_i}$  : Utility with policy  $\pi_i$


$U_i$ : Utilities at  $i$ -th iteration

$\pi_i$ : Policy recommended by value iteration if algorithm stops at  $i$ -th iteration

$U$ : Optimal Utility

# Policy Iteration

- Start with a random policy  $\pi_i$
- Policy Evaluation:
  - Calculate utility values for the policy  $\pi_i$
- Policy Improvement:
  - Calculate new policy  $\pi_{i+1}$  based on utility values
- Repeat the above steps till no change in policy



# Policy Iteration

**function** POLICY-ITERATION( $mdp$ ) **returns** a policy

**inputs:**  $mdp$ , an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$

**local variables:**  $U$ , a vector of utilities for states in  $S$ , initially zero

$\pi$ , a policy vector indexed by state, initially random

**repeat**

$U \leftarrow \text{POLICY-EVALUATION}(\pi, U, mdp)$   $\longrightarrow$  Calculate utility of each state with policy  $\pi$

$unchanged? \leftarrow \text{true}$

*U values cannot be used for policy update (why?)*

**for each** state  $s$  **in**  $S$  **do**

*So obtain Q values from U values*

$a^* \leftarrow \underset{a \in A(s)}{\operatorname{argmax}} \text{Q-VALUE}(mdp, s, a, U)$   $\longrightarrow$  Update the policy  $\pi$  based on utilities

**if**  $\text{Q-VALUE}(mdp, s, a^*, U) > \text{Q-VALUE}(mdp, s, \pi[s], U)$  **then**

$\pi[s] \leftarrow a^*;$   $unchanged? \leftarrow \text{false}$   $\longrightarrow$  If Q values of new policy are strictly larger than old policy, update the policy

**until**  $unchanged?$

**return**  $\pi$

**function**  $Q\text{-VALUE}(mdp, s, a, U)$  **returns** a utility value  
return  $\sum_{s'} P(s'|a, s_1)[R(s, a, s') + \gamma U(s')]$

# Conclusion

- Markov Decision Processes
  - Models uncertainty in environment
  - Utility-based reasoning
- Inferencing
  - Value Iteration
  - Policy iteration

# Expected Utility of State $s_0$ with a Policy $\pi$

$$\begin{aligned} U^\pi(s_0) &= E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(s_t), s_{t+1}) \mid s_0, \pi(s_0) \right] \\ &= E \left[ R(s_0, \pi(s_0), s_1) + \sum_{t=1}^{\infty} \gamma^t R(s_t, \pi(s_t), s_{t+1}) \mid s_0, \pi(s_0) \right] \\ &= \sum_{s_1} P(s_1 \mid s_0, \pi(s_0)) \left[ R(s_0, \pi(s_0), s_1) + E \left[ \sum_{t=1}^{\infty} \gamma^t R(s_t, \pi(s_t), s_{t+1}) \mid s_1, \pi(s_1) \right] \right] \\ &= \sum_{s_1} P(s_1 \mid s_0, \pi(s_0)) \left[ R(s_0, \pi(s_0), s_1) + \gamma E \left[ \sum_{t=1}^{\infty} \gamma^{t-1} R(s_t, \pi(s_t), s_{t+1}) \mid s_1, \pi(s_1) \right] \right] \\ &= \sum_{s_1} P(s_1 \mid s_0, \pi(s_0)) [R(s_0, \pi(s_0), s_1) + \gamma U^\pi(s_1)] \end{aligned}$$

$$P(s_0, \pi(s_0), s_1, \pi(s_1), \dots) = P(s_0) P(s_1 \mid s_0, \pi(s_0)) P(s_2 \mid s_1, \pi(s_1)) P(s_3 \mid s_2, \pi(s_2)) \dots$$