



# IT5005 Artificial Intelligence

Sirigina Rajendra Prasad  
AY2025/2026: Semester 1

## Tutorial 9: Transformers

# Data Preparation: Corpus and Dictionary

- Dictionary class
  - Bidirectional mapping from token ID to token
- Corpus class
  - Converts given text data into a sequence of tokens

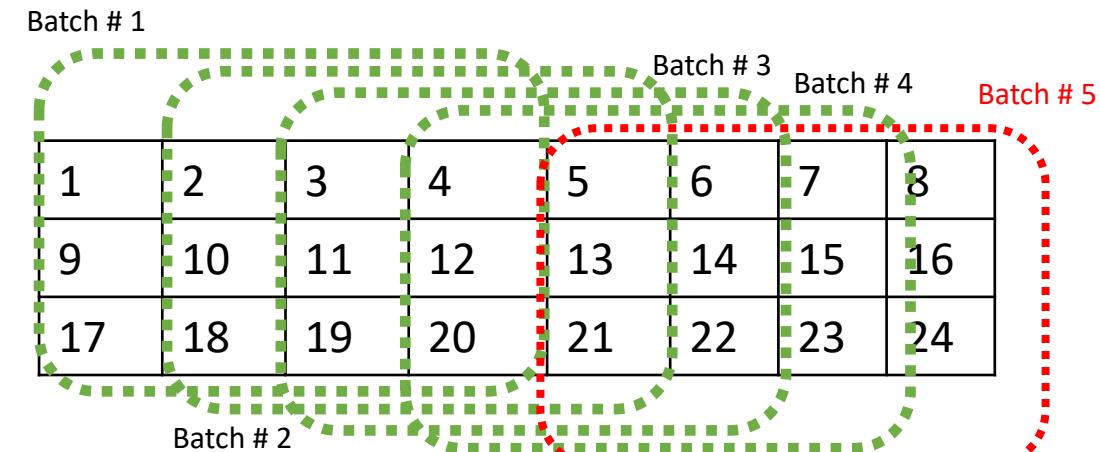
# Data Preparation

- Data is loaded in batches
- Two parameters to define a batch: *batch\_size* and *sequence\_length*
  - *batch\_size* is the number of parallel sequences in a batch that are processed simultaneously
  - *sequence\_length* is the length of the sequence in a batch
- Each batch is of dimension (*batch\_size*, *sequence\_length*)

# Data Preparation

- $data\_size$  = Total available tokens
- $batch\_size$  = Number of parallel sequences
- $tokens\_per\_sequence = \frac{data\_size}{batch\_size}$  -> integer division. (`_prepare_data` method)
- $seq\_len$  = window size
- $num\_batches = tokens\_per\_sequence - seq\_length$
- $start\_idx$  = starting position of the window
- Constraints
  - $seq\_len \leq tokens\_per\_sequence$
  - $start\_idx \leq tokens\_per\_sequence - seq\_len - 1$

# Data Preparation



- Example:
  - `corpus.data = [1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20,21,22,23,24,25,26]`
  - $seq\_len = 4$  (window size = 4 for sliding window)
  - $num\_batches = tokens\_per\_sequence - seq\_length = 8 - 4 = 4$
- We cannot use Batch #5. Why?
  - No target for Token 24.

# Preparing Data for Training

[1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20,21,22,23,24,25,26]

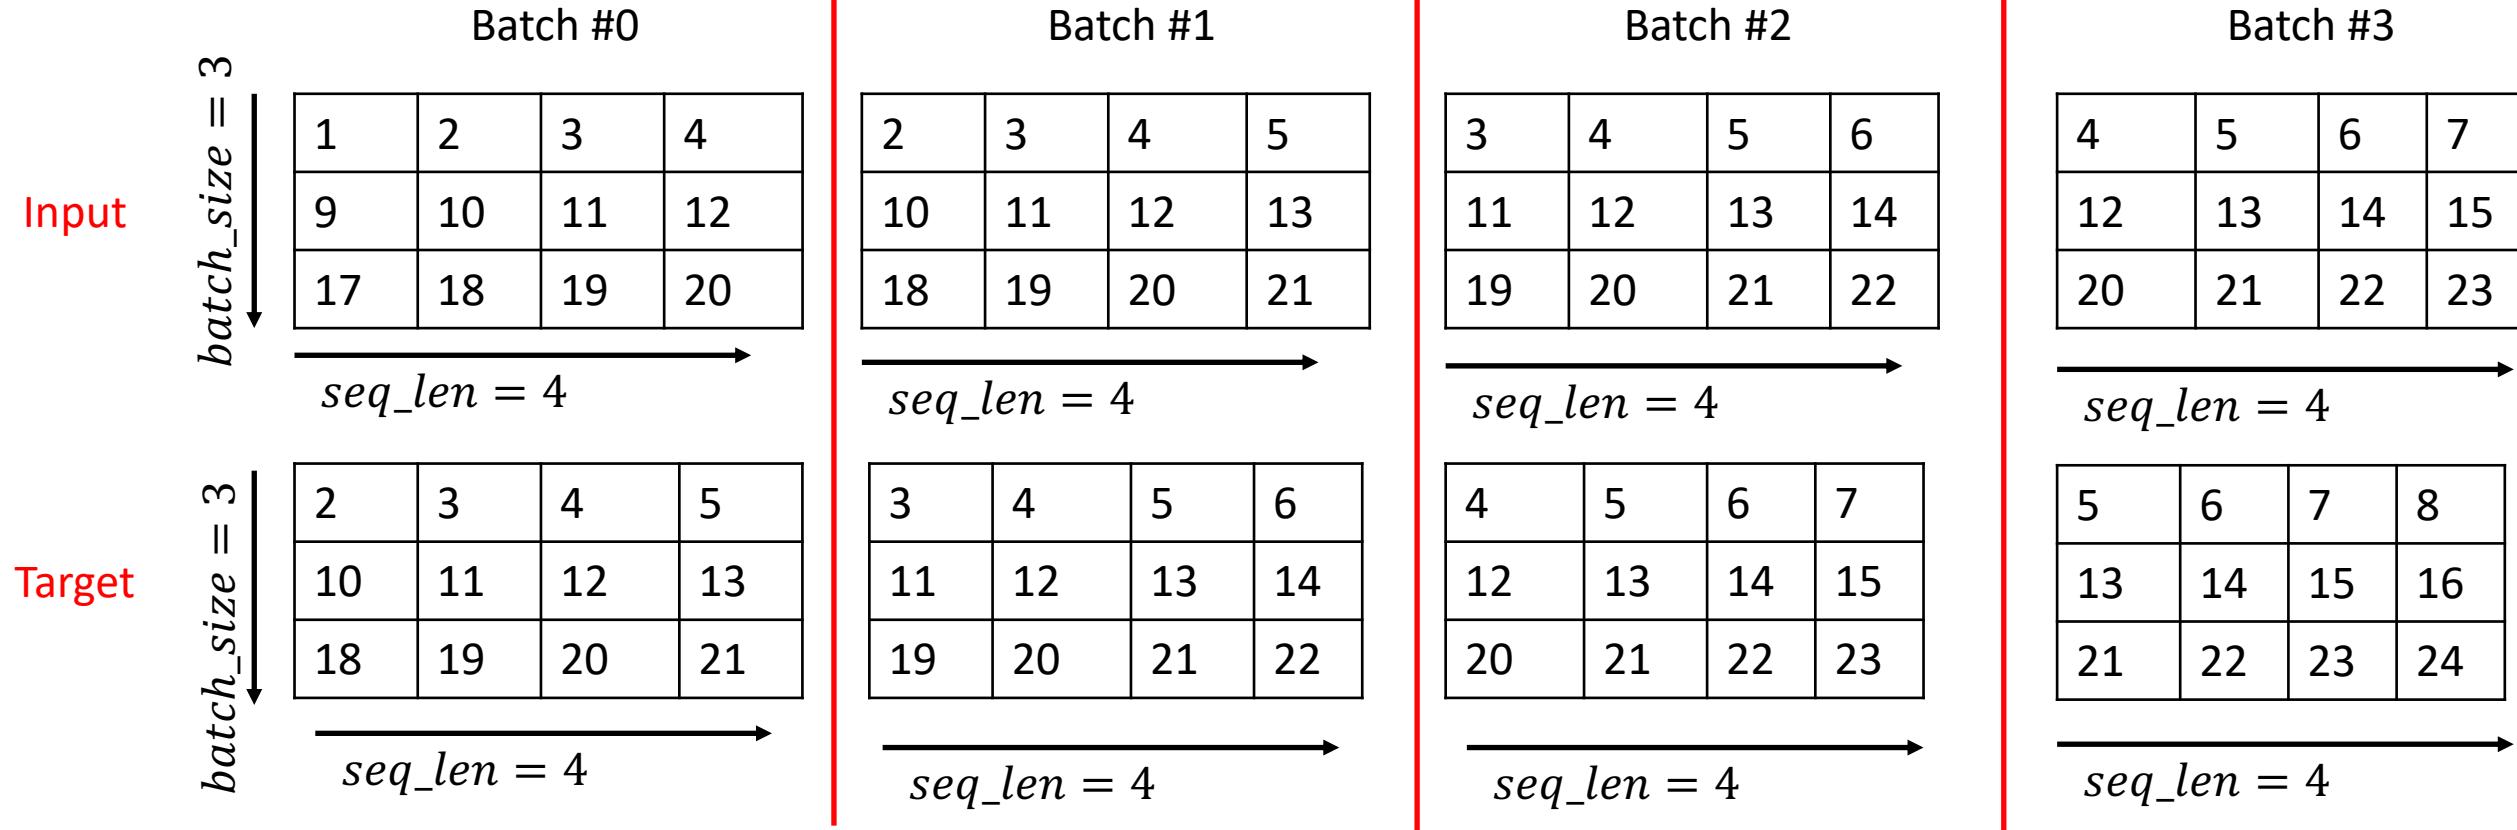
$$tokens\_per\_sequence = \frac{data\_size}{batch\_size} = 26/3 = 8$$

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22	23	24

# Batchification

Input Data

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22	23	24



`get_batch(i)` would pick  $i$  -th batch

# Batch #0

Input

$$\begin{array}{c} \xrightarrow{batch\_size = 3} \\ \xrightarrow{seq\_len = 4} \end{array}$$

1	2	3	4
9	10	11	12
17	18	19	20

Target

$$\begin{array}{c} \xrightarrow{batch\_size = 3} \\ \xrightarrow{seq\_len = 4} \end{array}$$

2	3	4	5
10	11	12	13
18	19	20	21

Input	Target
1	2
1,2	3
1,2,3	4
1,2,3,4	5
9	10
9,10	11
9,10,11	12
9,10,11,12	13
17	18
17,18	19
17,18,19	20
17,18,19,20	21

# Batch #1

Input

2	3	4	5
10	11	12	13
18	19	20	21

Target

$batch\_size = 3$  ↓

3	4	5	6
11	12	13	14
19	20	21	22

seq\_len = 4 →

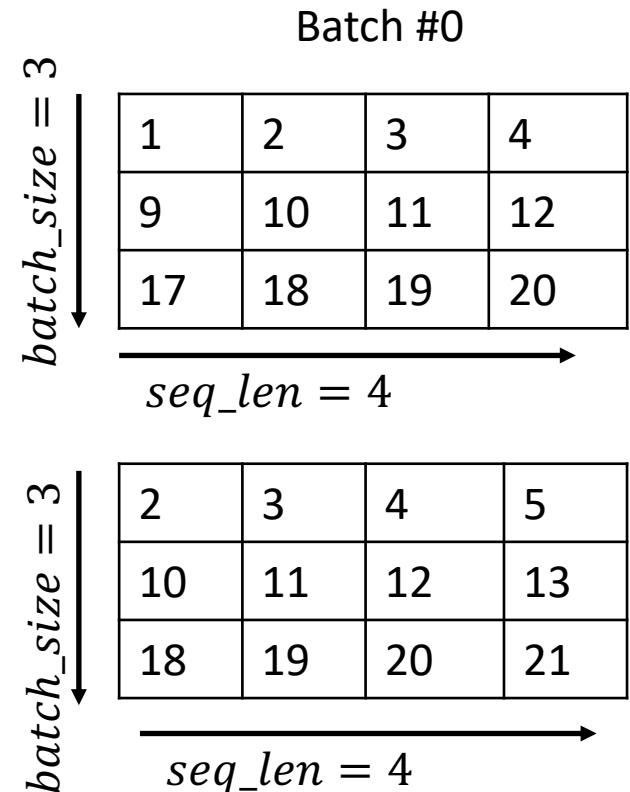
Input	Target
2	3
2,3	4
2,3,4	5
2,3,4,5	6
10	11
10,11	12
10,11,12	13
10,11,12,13	14
18	19
18,19	20
18,19,20	21
18,19,20,21	22

# How to choose batches?

- Larger batches
  - More stable gradients, smoother convergence
- Smaller batches
  - Noisier gradients, can help escape local minima
- Batch size
  - 32-128 for most language models
- Learning Rate Scaling:
  - Larger batches often require higher learning rates (why?)
  - Scale LR proportionally to batch size

# How to choose sequence length?

- Longer sequences
  - provide more training data
  - Better long-range dependencies, more context
- Shorter sequences
  - Faster training, less memory
- Trade-off
  - Computational cost grows quadratically with sequence length in attention



# Why Learning Rate Scheduler?

- Need to control the learning rate dynamically during training.
  - Improve convergence.
  - Adapt to different training phases (warmup, decay).
- Common Schedulers
  - Linear Decay
  - Cosine Decay

# Warmup + Decay Strategy

- Warmup
  - Start with a small LR, gradually increase.
- Decay
  - Reduce LR after warmup to fine-tune learning.
- Why warmup + Decay?
  - Stabilizes training in early steps.
  - Prevents gradient explosion.
  - Helps large models converge smoothly.

# Linear vs Cosine Decay

- Warmup

- $lr(t) = lr_{base} \frac{t}{W}$ .  $0 \leq t < W$

- **total\_steps**: Total training steps
- **warmup\_ratio**: Fraction of steps for warmup.
- **base\_lr**: Starting learning rate.
- **min\_lr**: Floor value to prevent LR from going too low.

- Linear Decay

- $lr(t) = \max\left(lr_{min}, lr_{base} \left(1 - \frac{t-W}{T-W}\right)\right) \quad W \leq t \leq T$

$$W = \text{warmup\_ratio} \times \text{total\_steps}$$

- LR decreases linearly after warmup.
- Simple and predictable.

- Cosine Decay

- $lr(t) = \max\left(lr_{min}, lr_{base} \frac{1}{2} \left(1 + \cos\left(2\pi \text{cycles} \frac{t-W}{T-W}\right)\right)\right) \quad W \leq t \leq T$

- LR follows a cosine curve.
- Avoids local minima

# Q3: Layer Norm

Let  $y$  be the output of Post-LN and it can be written as

$$y = \text{LN}(\mathbf{x} + \text{FFN}(\mathbf{x})).$$

Assume three input tokens with the corresponding embedding vectors  $\mathbf{x}_i, i \in \{1, 2, 3\}$ :

$$\mathbf{x}_1 = \begin{bmatrix} 1 \\ 0 \\ 2 \end{bmatrix}, \quad \mathbf{x}_2 = \begin{bmatrix} 3 \\ 1 \\ 4 \end{bmatrix}, \quad \mathbf{x}_3 = \begin{bmatrix} 5 \\ 2 \\ 6 \end{bmatrix}$$

FFN is a two-layer MLP with ReLU:

$$\text{FFN}(x) = W_2^T \text{ReLU}(W_1^T x + b_1) + b_2,$$

with

$$W_1 = I_3, \quad b_1 = \begin{bmatrix} 0.5 \\ -0.5 \\ 0 \end{bmatrix}, \quad W_2 = I_3, \quad b_2 = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}.$$

LayerNorm parameters (per feature):

$$\gamma = \begin{bmatrix} 2 \\ 1 \\ 0.5 \end{bmatrix}, \quad \beta = \begin{bmatrix} 0.5 \\ -0.5 \\ 1.0 \end{bmatrix}$$

Compute the output of Layer Normalization during training and inference.

# Layer Normalization

Minibatch:  $n = 5$  samples and  $m = 4$  features

- For each sample with  $m$  features:
  - Calculate **per-sample** mean and variance:
    - $\mu_i = \frac{1}{m} \sum_{i=1}^m x_i^{(j)}, i = \{1, 2, \dots, n\}$
    - $\sigma_i^2 = \frac{1}{m} \sum_{i=1}^m (x_i^{(j)} - \mu_i)^2, i = \{1, 2, \dots, n\}$
  - Normalize the data
    - $\bar{x}_i = \frac{x_i - \mu_i}{\sqrt{\sigma_i^2 + \epsilon}}$ , where  $\epsilon$  is a small constant for numerical stability
  - Scale and shift with learnable parameters  $\gamma \in \mathbf{R}^{m \times 1}$  and  $\beta \in \mathbf{R}^{m \times 1}$ 
    - $y_i = \gamma \circ \bar{x}_i + \beta$
  - Learnable parameters allows undoing of normalization if needed

	$x_1$	$x_2$	$x_3$	$x_4$	$\mu_1$	$\sigma_1^2$
$x_i^{(1)}$	2	80	400	0.5	$\mu_2$	$\sigma_2^2$
$x_i^{(2)}$	4	90	300	0.7	$\mu_3$	$\sigma_3^2$
$x_i^{(3)}$	6	70	500	0.4	$\mu_4$	$\sigma_4^2$
$x_i^{(4)}$	8	85	600	0.6	$\mu_5$	$\sigma_5^2$
$x_i^{(5)}$	10	95	200	0.8		

$$x_i = \begin{bmatrix} x_i^{(1)} \\ \vdots \\ x_i^{(n)} \end{bmatrix} \in \mathbf{R}^{n \times 1}$$

$$\bar{x}_i \in \mathbf{R}^{n \times 1}$$

# Layer Normalization

Let  $y$  be the output of Post-LN and it can be written as

$$y = \text{LN}(\mathbf{x} + \text{FFN}(\mathbf{x})).$$

Assume three input tokens with the corresponding embedding vectors  $\mathbf{x}_i, i \in \{1, 2, 3\}$ :

$$\mathbf{x}_1 = \begin{bmatrix} 1 \\ 0 \\ 2 \end{bmatrix}, \quad \mathbf{x}_2 = \begin{bmatrix} 3 \\ 1 \\ 4 \end{bmatrix}, \quad \mathbf{x}_3 = \begin{bmatrix} 5 \\ 2 \\ 6 \end{bmatrix}$$

FFN is a two-layer MLP with ReLU:

$$\text{FFN}(x) = W_2^T \text{ReLU}(W_1^T x + b_1) + b_2,$$

with

$$W_1 = I_3, \quad b_1 = \begin{bmatrix} 0.5 \\ -0.5 \\ 0 \end{bmatrix}, \quad W_2 = I_3, \quad b_2 = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}.$$

LayerNorm parameters (per feature):

$$\gamma = \begin{bmatrix} 2 \\ 1 \\ 0.5 \end{bmatrix}, \quad \beta = \begin{bmatrix} 0.5 \\ -0.5 \\ 1.0 \end{bmatrix}$$

Compute the output of Layer Normalization during training and inference.

$$\mathbf{x}_1 + \text{FFN}(\mathbf{x}_1) = \begin{bmatrix} 2.5 \\ 0 \\ 4 \end{bmatrix}$$

$$\mu_1 = 2.166$$

$$\sigma_1 = 1.649$$

$$\frac{\mathbf{x}_1 - \mu_1}{\sigma_1} = \begin{bmatrix} 0.202 \\ -1.31 \\ 1.11 \end{bmatrix}$$

$$\text{LN}(\mathbf{x}_i + \text{FFN}(\mathbf{x}_i)) = \gamma \circ \left( \frac{\mathbf{x}_i - \mu_i}{\sigma_i} \right) + \beta$$

$$\text{LN}(\mathbf{x}_1 + \text{FFN}(\mathbf{x}_1)) = \begin{bmatrix} 2 \\ 1 \\ 0.5 \end{bmatrix} \circ \left( \frac{\mathbf{x}_1 - \mu_1}{\sigma_1} \right) + \begin{bmatrix} 0.5 \\ -0.5 \\ 1.0 \end{bmatrix} = \begin{bmatrix} 0.904 \\ -1.813 \\ 1.555 \end{bmatrix}$$

# Layer Normalization

Let  $y$  be the output of Post-LN and it can be written as

$$y = \text{LN}(\mathbf{x} + \text{FFN}(\mathbf{x})).$$

Assume three input tokens with the corresponding embedding vectors  $\mathbf{x}_i, i \in \{1, 2, 3\}$ :

$$\mathbf{x}_1 = \begin{bmatrix} 1 \\ 0 \\ 2 \end{bmatrix}, \quad \mathbf{x}_2 = \begin{bmatrix} 3 \\ 1 \\ 4 \end{bmatrix}, \quad \mathbf{x}_3 = \begin{bmatrix} 5 \\ 2 \\ 6 \end{bmatrix}$$

FFN is a two-layer MLP with ReLU:

$$\text{FFN}(x) = W_2^T \text{ReLU}(W_1^T x + b_1) + b_2,$$

with

$$W_1 = I_3, \quad b_1 = \begin{bmatrix} 0.5 \\ -0.5 \\ 0 \end{bmatrix}, \quad W_2 = I_3, \quad b_2 = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}.$$

LayerNorm parameters (per feature):

$$\gamma = \begin{bmatrix} 2 \\ 1 \\ 0.5 \end{bmatrix}, \quad \beta = \begin{bmatrix} 0.5 \\ -0.5 \\ 1.0 \end{bmatrix}$$

Compute the output of Layer Normalization during training and inference.

$$\mathbf{x}_2 + \text{FFN}(\mathbf{x}_2) = \begin{bmatrix} 6.5 \\ 1.5 \\ 8 \end{bmatrix}$$

$$\mu_2 = 5.33$$

$$\sigma_2 = 2.77$$

$$\frac{\mathbf{x}_2 - \mu_2}{\sigma_2} = \begin{bmatrix} 0.4198 \\ -1.379 \\ 0.959 \end{bmatrix}$$

$$\text{LN}(\mathbf{x}_i + \text{FFN}(\mathbf{x}_i)) = \gamma \circ \left( \frac{\mathbf{x}_i - \mu_i}{\sigma_i} \right) + \beta$$

$$\text{LN}(\mathbf{x}_2 + \text{FFN}(\mathbf{x}_2)) = \begin{bmatrix} 2 \\ 1 \\ 0.5 \end{bmatrix} \circ \left( \frac{\mathbf{x}_2 - \mu_2}{\sigma_2} \right) + \begin{bmatrix} 0.5 \\ -0.5 \\ 1.0 \end{bmatrix} = \begin{bmatrix} 1.339 \\ -1.879 \\ 1.479 \end{bmatrix}$$

# Layer Normalization

Let  $y$  be the output of Post-LN and it can be written as

$$y = \text{LN}(\mathbf{x} + \text{FFN}(\mathbf{x})).$$

Assume three input tokens with the corresponding embedding vectors  $\mathbf{x}_i, i \in \{1, 2, 3\}$ :

$$\mathbf{x}_1 = \begin{bmatrix} 1 \\ 0 \\ 2 \end{bmatrix}, \quad \mathbf{x}_2 = \begin{bmatrix} 3 \\ 1 \\ 4 \end{bmatrix}, \quad \mathbf{x}_3 = \begin{bmatrix} 5 \\ 2 \\ 6 \end{bmatrix}$$

FFN is a two-layer MLP with ReLU:

$$\text{FFN}(x) = W_2^T \text{ReLU}(W_1^T x + b_1) + b_2,$$

with

$$W_1 = I_3, \quad b_1 = \begin{bmatrix} 0.5 \\ -0.5 \\ 0 \end{bmatrix}, \quad W_2 = I_3, \quad b_2 = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}.$$

LayerNorm parameters (per feature):

$$\gamma = \begin{bmatrix} 2 \\ 1 \\ 0.5 \end{bmatrix}, \quad \beta = \begin{bmatrix} 0.5 \\ -0.5 \\ 1.0 \end{bmatrix}$$

Compute the output of Layer Normalization during training and inference.

$$\mathbf{x}_3 + \text{FFN}(\mathbf{x}_3) = \begin{bmatrix} 10.5 \\ 3.5 \\ 12 \end{bmatrix}$$

$$\mu_3 = 8.66$$

$$\sigma_3 = 3.70$$

$$\frac{\mathbf{x}_3 - \mu_3}{\sigma_3} = \begin{bmatrix} 0.4949 \\ -1.394 \\ 0.899 \end{bmatrix}$$

$$\text{LN}(\mathbf{x}_i + \text{FFN}(\mathbf{x}_i)) = \gamma \circ \left( \frac{\mathbf{x}_i - \mu_i}{\sigma_i} \right) + \beta$$

$$\text{LN}(\mathbf{x}_3 + \text{FFN}(\mathbf{x}_3)) = \begin{bmatrix} 2 \\ 1 \\ 0.5 \end{bmatrix} \circ \left( \frac{\mathbf{x}_3 - \mu_3}{\sigma_3} \right) + \begin{bmatrix} 0.5 \\ -0.5 \\ 1.0 \end{bmatrix} = \begin{bmatrix} 1.489 \\ -1.894 \\ 1.449 \end{bmatrix}$$

# Why Pre-LN preferred over Post-LN?

---

## **On Layer Normalization in the Transformer Architecture**

---

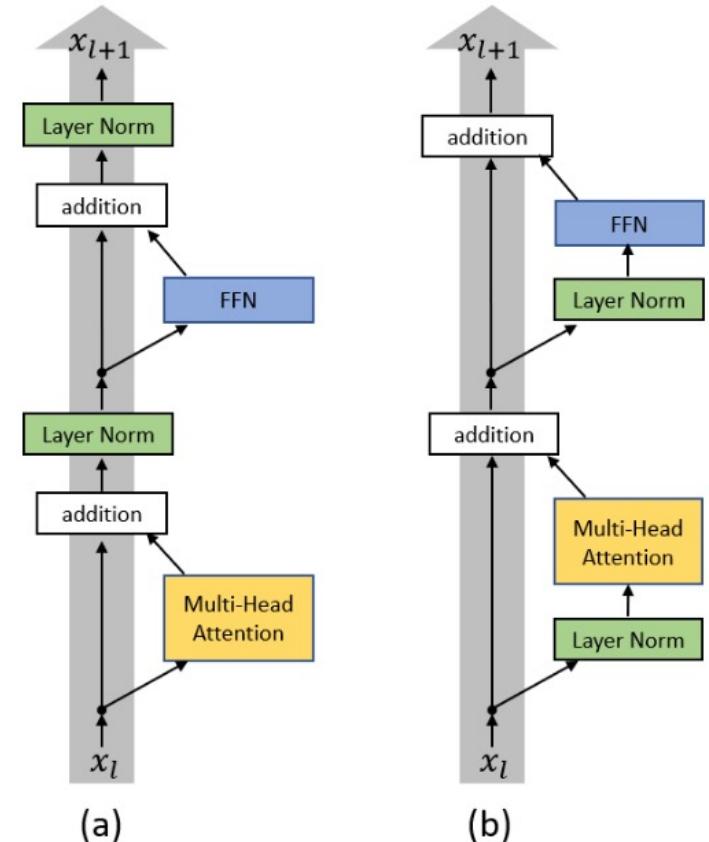
**Ruixin Xiong<sup>†\*</sup> <sup>1,2</sup> Yunchang Yang<sup>\*</sup> <sup>3</sup> Di He<sup>4,5</sup> Kai Zheng<sup>4</sup> Shuxin Zheng<sup>5</sup> Chen Xing<sup>6</sup> Huishuai Zhang<sup>5</sup>**  
**Yanyan Lan<sup>1,2</sup> Liwei Wang<sup>4,3</sup> Tie-Yan Liu<sup>5</sup>**

<https://arxiv.org/pdf/2002.04745>

# Post-LN

```
class TransformerBlock(nn.Module):
    def __init__(self, d_model, num_heads, d_ff, dropout=0.1):
        super().__init__()
        self.attention = MultiHeadAttention(d_model, num_heads)
        self.feed_forward = PositionwiseFeedForward(d_model, d_ff)
        self.norm1 = nn.LayerNorm(d_model)
        self.norm2 = nn.LayerNorm(d_model)
        self.dropout = nn.Dropout(dropout)

    def forward(self, x, mask=None):
        attn_out = self.attention(x, x, x, mask)
        x = self.norm1(x + self.dropout(attn_out))
        ff_out = self.feed_forward(x)
        x = self.norm2(x + self.dropout(ff_out))
        return x
```



(a) Post-LN Transformer layer; (b) Pre-LN Transformer

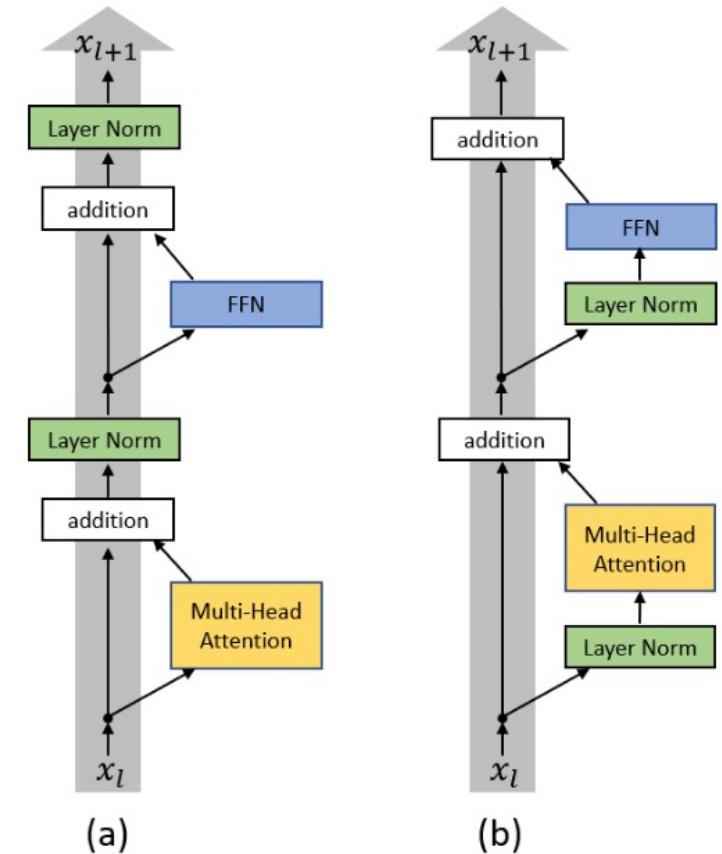
# Pre-LN

```
class TransformerBlock(nn.Module):
    def __init__(self, d_model, num_heads, d_ff, dropout=0.1):
        super().__init__()
        self.attention = MultiHeadAttention(d_model, num_heads)
        self.feed_forward = PositionwiseFeedForward(d_model, d_ff)
        self.norm1 = nn.LayerNorm(d_model)
        self.norm2 = nn.LayerNorm(d_model)
        self.dropout = nn.Dropout(dropout)

    def forward(self, x, mask=None):
        # Pre-LN attention: normalize input, apply attention, then residual add
        x_norm = self.norm1(x)
        attn_out = self.attention(x_norm, x_norm, x_norm, mask)
        x = x + self.dropout(attn_out)

        # Pre-LN FFN: normalize current x, apply FFN, then residual add
        x_norm = self.norm2(x)
        ff_out = self.feed_forward(x_norm)
        x = x + self.dropout(ff_out)

    return x
```



(a) Post-LN Transformer layer; (b) Pre-LN Transformer

# Why Pre-LN preferred over Post-LN?

- Post-LN block

- Attention sublayer:
    - $y = \text{LN}(x + \text{Attn}(x))$
  - FFN sublayer:
    - $y = \text{LN}(x + \text{FFN}(x))$

- Pre-LN block

- Attention sublayer:
    - $y = x + (\text{Attn}(\text{LN}(x))$
  - FFN sublayer:
    - $y = x + (\text{Attn}(\text{LN}(x))$

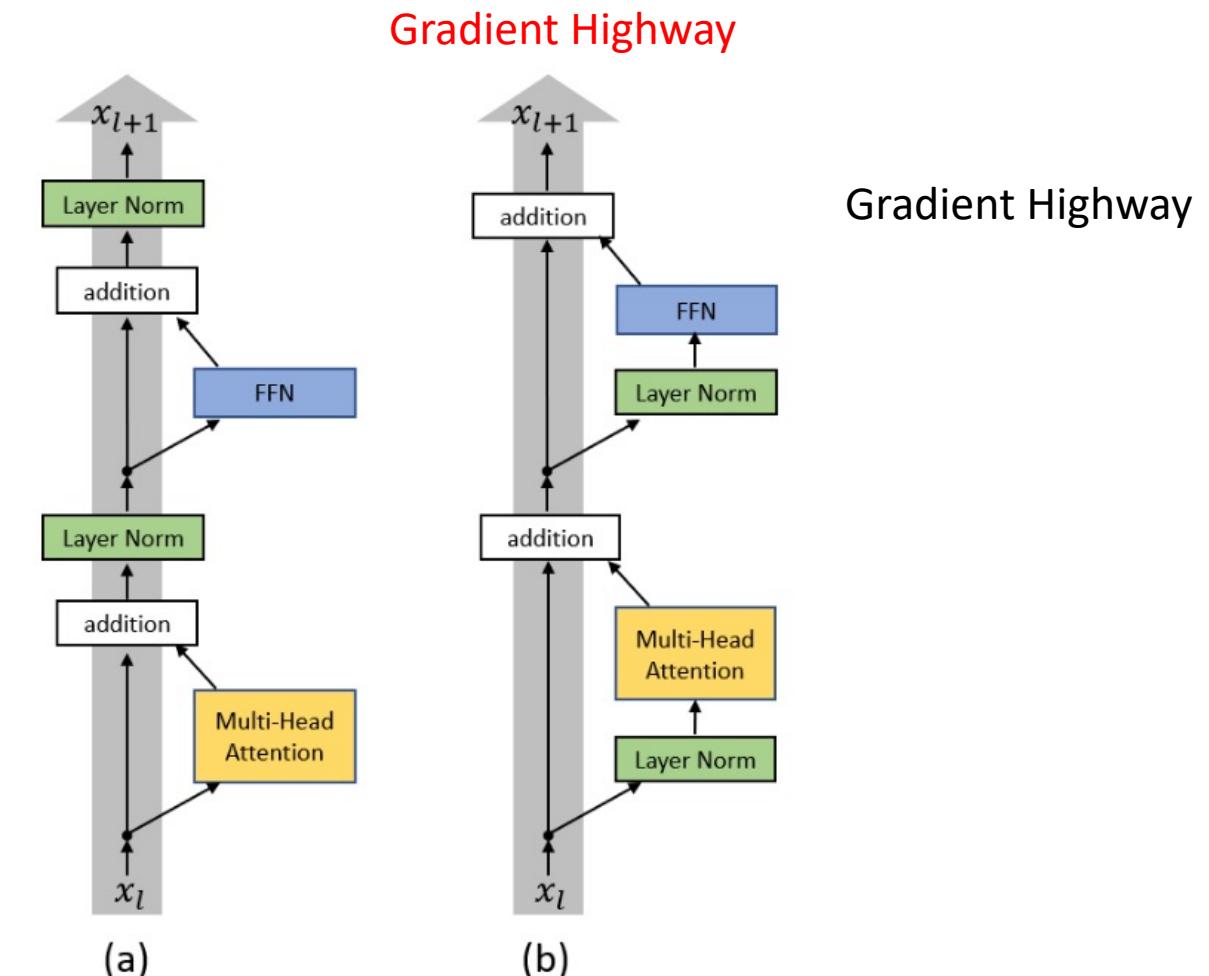


Figure 1. (a) Post-LN Transformer layer; (b) Pre-LN Transformer layer.

# Why Pre-LN preferred over Post-LN?

$$\frac{\partial y}{\partial x} = \frac{\partial \text{LN}}{\partial x} \left( I + \frac{\partial \text{Att}}{\partial x} \right)$$

$$\frac{\partial y}{\partial x} = \frac{\partial \text{LN}}{\partial x} \left( I + \frac{\partial \text{FFN}}{\partial x} \right)$$

Gradients via residual paths are multiplied by gradients of LN block

It can shrink or distort the gradients

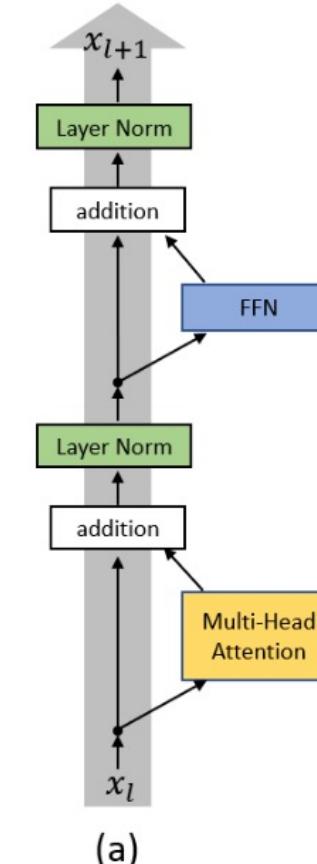
$$\frac{\partial y}{\partial x} = I + \frac{\partial \text{Attn}}{\partial \text{LN}} \frac{\partial \text{LN}}{\partial x}$$

$$\frac{\partial y}{\partial x} = I + \frac{\partial \text{FFN}}{\partial \text{LN}} \frac{\partial \text{LN}}{\partial x}$$

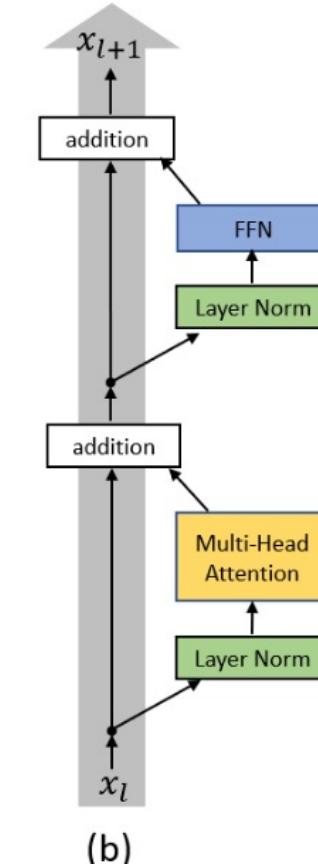
Clean 'identity' path for gradient or gradient highway via residual paths

Converges robustly and allows higher learning rates

Gradient Highway



(a)



(b)

Figure 1. (a) Post-LN Transformer layer; (b) Pre-LN Transformer layer.

# Pre-LN

- Learning Rate scheduler
  - Can use larger learning rate compared to post-LN
  - No warmup needed
    - 'warmup\_ratio': 0.0,

# How to improve performance?

- Large  $d_{model}$ . (start with small 128 and increase it to 256)
- Increase sequence length (start with 100 and increase it to 256)
  - Larger context
- Different activation (start with ReLU and change it to GELU)
- Pre-LN
- Lower drop out
- Stop the training once the validation loss plateaus
- If validation loss increases after above changes, revert the change...it's a sign of overfitting